

Analysis: Classification Trees & Random Forest

Hannah Norman

```
library(caret)

## Loading required package: ggplot2
## Loading required package: lattice
library(mlbench)
library(rattle)

## Loading required package: tibble
## Loading required package: bitops
## Rattle: A free graphical interface for data science with R.
## Version 5.5.1 Copyright (c) 2006-2021 Togaware Pty Ltd.
## Type 'rattle()' to shake, rattle, and roll your data.
library(randomForest)

## randomForest 4.7-1.1
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
## The following object is masked from 'package:rattle':
##
##     importance
## The following object is masked from 'package:ggplot2':
##
##     margin
library(ROCR)
```

Original NASA Asteroids Data Set

```
nasa_orig <- read.csv("nasa_original.csv")
# nasa_orig
```

Cleaned NASA Asteroids Data Set

```
nasa <- read.csv("nasa.csv")
# nasa
```

Numerical and Graphical Summaries of Response Variable

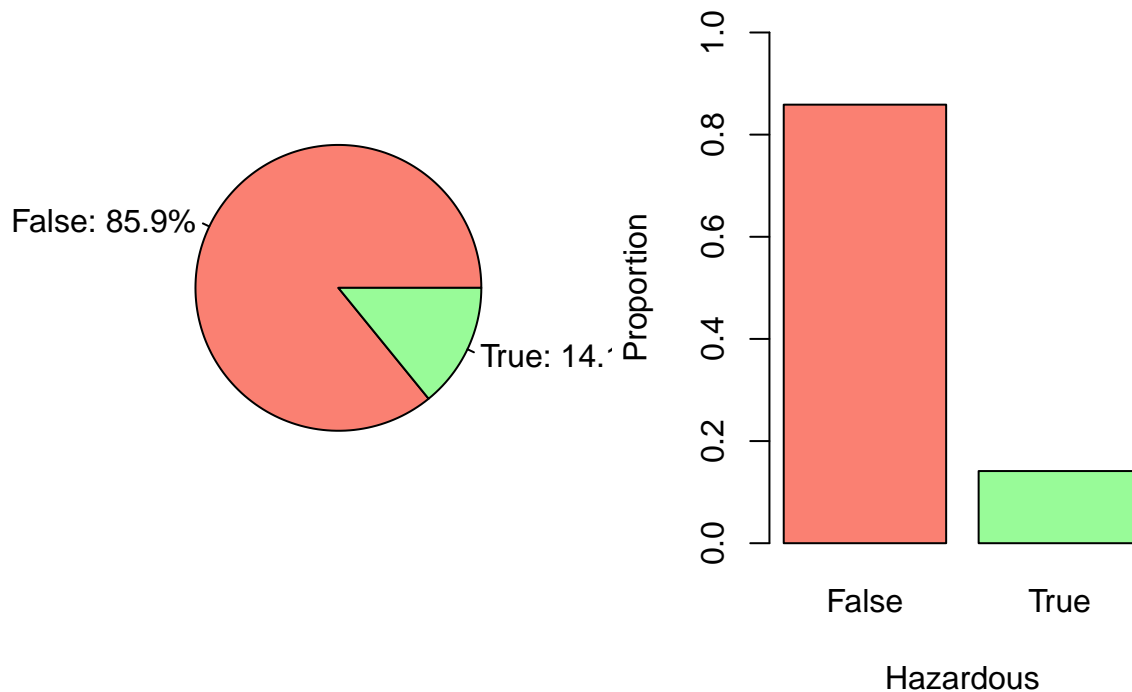
```
prop.hazardous <- prop.table(table(nasa$Hazardous))
prop.hazardous

##
##      False      True
## 0.8587204 0.1412796

par(mfrow=c(1,2))

# pie chart
count.hazardous <- table(nasa$Hazardous)
lbls <- paste(levels(as.factor(nasa$Hazardous)), ":", " ",
              round(prop.hazardous,3)*100, "%", sep="")
pie(count.hazardous, labels=lbls, col=c("salmon", "palegreen"))

# bar plot
barplot(prop.hazardous, xlab="Hazardous", ylab="Proportion", ylim=c(0, 1.0),
        col=c("salmon", "palegreen"))
```



Analyses of Predictors

(in search of outliers that might skew analysis)

```
# numerical summary + box plots
summary(nasa)
```

```
## Absolute.Magnitude Est.Dia.in.KM.min. Est.Dia.in.KM.max. Est.Dia.in.KM.range
## Min. :14.40 Min. :0.001011 Min. :0.00226 Min. :0.001249
## 1st Qu.:20.40 1st Qu.:0.030518 1st Qu.:0.06824 1st Qu.:0.037722
## Median :22.30 Median :0.092163 Median :0.20608 Median :0.113919
## Mean :22.53 Mean :0.172936 Mean :0.38670 Mean :0.213761
```

```
## 3rd Qu.:24.70      3rd Qu.:0.221083      3rd Qu.:0.49436      3rd Qu.:0.273273
## Max. :32.10      Max. :3.503926      Max. :7.83502      Max. :4.331091
## Close.Approach.Date Relative.Velocity.in.KM.per.sec Miss.Dist.in.KM
## Min. :19950101      Min. : 0.8002      Min. : 26610
## 1st Qu.:20010765      1st Qu.: 8.1683      1st Qu.:16224266
## Median :20070908      Median :12.3900      Median :37032388
## Mean :20066238      Mean :13.6152      Mean :36468199
## 3rd Qu.:20120922      3rd Qu.:17.5084      3rd Qu.:56281652
## Max. :20160908      Max. :43.7899      Max. :74781600
## Orbit.Uncertainty Minimum.Orbit.Intersection Jupiter.Tisserand.Invariant
## Min. :0.000      Min. :0.0000021      Min. :2.196
## 1st Qu.:1.000      1st Qu.:0.0151384      1st Qu.:3.804
## Median :5.000      Median :0.0473248      Median :4.798
## Mean :4.099      Mean :0.0823078      Mean :4.852
## 3rd Qu.:7.000      3rd Qu.:0.1248985      3rd Qu.:5.774
## Max. :9.000      Max. :0.4778910      Max. :9.025
## Eccentricity Semi.Major.Axis Inclination Asc.Node.Longitude
## Min. :0.01296      Min. :0.6159      Min. : 0.01451      Min. : 0.0019
## 1st Qu.:0.24918      1st Qu.:1.0530      1st Qu.: 4.78919      1st Qu.: 83.6762
## Median :0.38253      Median :1.3349      Median : 9.68518      Median :173.6808
## Mean :0.39329      Mean :1.4850      Mean :12.83906      Mean :173.5626
## 3rd Qu.:0.52595      3rd Qu.:1.8388      3rd Qu.:18.38036      3rd Qu.:258.6316
## Max. :0.96026      Max. :3.9908      Max. :75.40667      Max. :359.9059
## Orbital.Period Perihelion.Distance Perihelion.Arg Aphelion.Dist
## Min. : 176.6      Min. :0.08074      Min. : 0.0069      Min. :0.8038
## 1st Qu.: 394.7      1st Qu.:0.67808      1st Qu.: 95.6430      1st Qu.:1.3191
## Median : 563.3      Median :0.87390      Median :188.5239      Median :1.7918
## Mean : 692.9      Mean :0.84320      Mean :184.1443      Mean :2.1268
## 3rd Qu.: 910.7      3rd Qu.:1.01935      3rd Qu.:272.5059      3rd Qu.:2.7545
## Max. :2912.0      Max. :1.29983      Max. :359.9931      Max. :6.8918
## Mean.Anomaly Mean.Motion Hazardous
## Min. : 0.0032      Min. :0.1236      Length:3079
## 1st Qu.: 83.5492      1st Qu.:0.3953      Class :character
## Median :183.9847      Median :0.6391      Mode :character
## Mean :180.1871      Mean :0.6810
## 3rd Qu.:276.3719      3rd Qu.:0.9121
## Max. :359.9180      Max. :2.0390
```


```
par(mfrow=c(2,4))
boxplot(nasa$Est.Dia.in.KM.range, main="Est.Dia.in.KM.range")
boxplot(nasa$Relative.Velocity, main="Relative.Velocity")
boxplot(nasa$Minimum.Orbit.Intersection, main="Minimum.Orbit.Intersection")
boxplot(nasa$Inclination, main="Inclination")
boxplot(nasa$Orbital.Period, main="Orbital.Period")
boxplot(nasa$Aphelion.Dist, main="Aphelion.Dist")
boxplot(nasa$Mean.Motion, main="Mean.Motion")
```

```
# to find row number of outlier observations
nasa[which.max(nasa$Est.Dia.in.KM.range),]
```


```
## Absolute.Magnitude Est.Dia.in.KM.min. Est.Dia.in.KM.max.
## 695 14.4 3.503926 7.835018
## Est.Dia.in.KM.range Close.Approach.Date Relative.Velocity.in.KM.per.sec
## 695 4.331091 20001222 21.19854
## Miss.Dist.in.KM Orbit.Uncertainty Minimum.Orbit.Intersection
```

```
## 695      21340634      0      0.0282524
##      Jupiter.Tisserand.Invariant Eccentricity Semi.Major.Axis Inclination
## 695      3.573      0.6343498      1.982214      6.705068
##      Asc.Node.Longitude Orbital.Period Perihelion.Distance Perihelion.Arg
## 695      294.8956      1019.352      0.7247968      236.3403
##      Aphelion.Dist Mean.Anomaly Mean.Motion Hazardous
## 695      3.239631      338.28      0.3531656      True
```


Est.Dia.in.KM.range




Relative.Velocity




Minimum.Orbit.Intersect




Inclination




Orbital.Period



Aphelion.Dist



Mean.Motion



Predictive Performance Stats

```
get_stats <- function(CM) {
  TP <- CM[2,2]
  FP <- CM[1,2]
  TN <- CM[1,1]
  FN <- CM[2,1]

  acc <- (TP+TN) / (TP+TN+FN+FP)
  err <- (FP+FN) / (TP+TN+FN+FP)
  pre <- (TP) / (TP+FP)
  sen <- (TP) / (TP+FN)
  spe <- (TN) / (TN+FP)
  fme <- (2*pre*sen) / (pre+sen)
  mcc_denom <- sqrt(TP+FP)*sqrt(TP+FN)*sqrt(TN+FP)*sqrt(TN+FN)
  mcc <- (TP*TN - FP*FN) / mcc_denom

  name <- c("accuracy", "error rate", "precision", "sensitivity", "specificity", "F-measure", "Matthew's")
  value <- c(acc, err, pre, sen, spe, fme, mcc)
  stats <- data.frame(name, value)

  return (stats)
}
```

Classification Tree Analysis

1. Make sure that the model assumptions, if any, are satisfied.

No model assumptions of decision trees to be satisfied? We are fitting a classification tree as opposed to a regression tree because the response variable is categorical and binary.

2. Assess the model fit and perform diagnostics, if appropriate.

Both methods appear to give identical results. Note the identical accuracy and kappa ratings across all cp tuning parameters in the plain rpart method...


```
set.seed(1)

# rpart
nasa.CVrpart <- train(Hazardous ~ ., data=nasa,
                      method="rpart",
                      tuneGrid = expand.grid(cp=seq(0.005, 0.05, length=10)),
                      trControl=trainControl(method="cv", number=10,
                                              savePredictions=TRUE,
                                              classProbs=TRUE,
                                              selectionFunction = "oneSE"))

nasa.CVrpart

## CART
##
## 3079 samples
## 20 predictor
## 2 classes: 'False', 'True'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 2771, 2771, 2772, 2770, 2771, 2772, ...
## Resampling results across tuning parameters:
##
##  cp      Accuracy  Kappa
##  0.005  0.993829  0.9743262
##  0.010  0.993829  0.9743262
##  0.015  0.993829  0.9743262
##  0.020  0.993829  0.9743262
##  0.025  0.993829  0.9743262
##  0.030  0.993829  0.9743262
##  0.035  0.993829  0.9743262
##  0.040  0.993829  0.9743262
##  0.045  0.993829  0.9743262
##  0.050  0.993829  0.9743262
##
## Accuracy was used to select the optimal model using the one SE rule.
## The final value used for the model was cp = 0.05.

# print tree
fancyRpartPlot(nasa.CVrpart$finalModel)
```



Rattle 2023-Apr-22 22:57:11 hannah

```


set.seed(1)

# rpart1SE
nasa.CVrpart1SE <- train(Hazardous ~ ., data=nasa,
                          method="rpart1SE",
                          trControl=trainControl(method="cv", number=10,
                                                  savePredictions=TRUE))

nasa.CVrpart1SE

## CART
##
## 3079 samples
## 20 predictor
## 2 classes: 'False', 'True'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 2771, 2771, 2772, 2770, 2771, 2772, ...
## Resampling results:
##
## Accuracy Kappa
## 0.993829 0.9743262

# print tree
fancyRpartPlot(nasa.CVrpart1SE$finalModel)
  
```



Rattle 2023-Apr-22 22:57:12 hannah

3. Identify tuning parameters to be used, if appropriate.

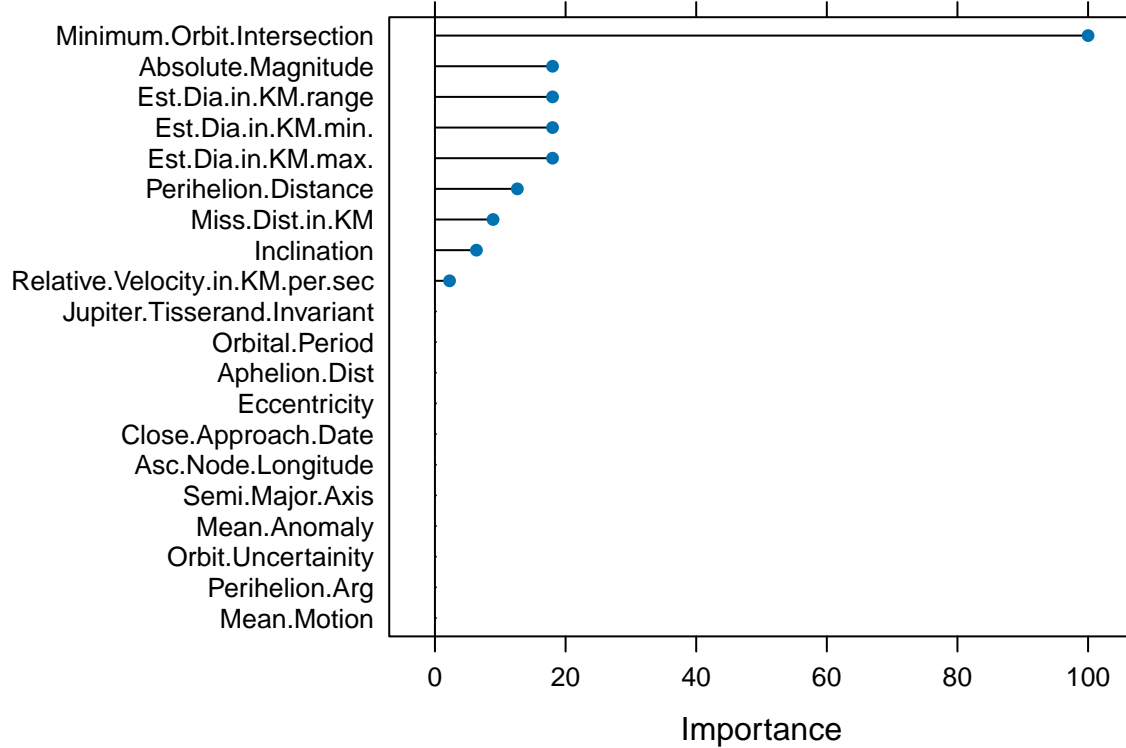
If the plain `rpart` method is utilized, it selects a `cp` (complexity parameter) value of 0.05 to maximize accuracy for the final model. Note that the accuracy and kappa values are identical for each of the `cp` values tried by the model, so any choice within that range should produce comparable results.

4. Identify and interpret the effect of selected variables.

```
# variable importance (rpart)
varImp(nasa.CVrpart)
```

```
## rpart variable importance
##
##                               Overall
## Minimum.Orbit.Intersection    100.000
## Absolute.Magnitude            18.004
## Est.Dia.in.KM.range           18.004
## Est.Dia.in.KM.min.            18.004
## Est.Dia.in.KM.max.            18.004
## Perihelion.Distance           12.614
## Miss.Dist.in.KM                8.894
## Inclination                    6.351
## Relative.Velocity.in.KM.per.sec 2.235
## Semi.Major.Axis                0.000
## Jupiter.Tisserand.Invariant    0.000
## Perihelion.Arg                 0.000
## Close.Approach.Date            0.000
## Aphelion.Dist                  0.000
## Mean.Motion                    0.000
## Mean.Anomaly                   0.000
## Asc.Node.Longitude             0.000
## Orbit.Uncertainty              0.000
## Orbital.Period                 0.000
## Eccentricity                   0.000
```

```
plot(varImp(nasa.CVrpart))
```




```
# variable importance (rpart1SE)
varImp(nasa.CVrpart1SE)
```

```
## rpart1SE variable importance
##
## Overall
## Minimum.Orbit.Intersection 100.000
## Absolute.Magnitude         18.004
## Est.Dia.in.KM.range        18.004
## Est.Dia.in.KM.min.         18.004
## Est.Dia.in.KM.max.         18.004
## Perihelion.Distance        12.614
## Miss.Dist.in.KM            8.894
## Inclination                 6.351
## Relative.Velocity.in.KM.per.sec 2.235
## Eccentricity                0.000
## Orbital.Period              0.000
## Mean.Motion                 0.000
## Semi.Major.Axis             0.000
## Perihelion.Arg              0.000
## Asc.Node.Longitude          0.000
## Aphelion.Dist               0.000
## Mean.Anomaly                0.000
## Close.Approach.Date         0.000
## Orbit.Uncertainty           0.000
## Jupiter.Tisserand.Invariant  0.000
```



```
plot(varImp(nasa.CVrpart1SE))
```




5. Evaluate the cross-validated (CV) predictive performance.

```
head(nasa.CVrpart$pred)
```

```
##      pred  obs rowIndex      False      True      cp Resample
## 1 False False      11 1.000000000 0.000000000 0.005 Fold01
## 2 False False      13 0.992572586 0.007427414 0.005 Fold01
## 3 False False      34 1.000000000 0.000000000 0.005 Fold01
## 4 False False      38 0.992572586 0.007427414 0.005 Fold01
## 5  True  True      66 0.002624672 0.997375328 0.005 Fold01
## 6  True  True      68 0.002624672 0.997375328 0.005 Fold01
```

```
pihatcv.rpart <- nasa.CVrpart$pred[nasa.CVrpart$pred$cp == 0.05,]
predcv.rpart <- prediction(pihatcv.rpart$True, pihatcv.rpart$obs)
perfcv.rpart <- performance(predcv.rpart, "tpr", "fpr")
plot(perfcv.rpart)
abline(a=0, b=1, lty=2)
```



```
aucCV.rpart <- performance(predcv.rpart, "auc")@y.values
aucCV.rpart
```

```
## [[1]]
## [1] 0.9884736
```

```
confMat <- table(pihatcv.rpart$obs, pihatcv.rpart$pred)
confMat
```

```
##
##      False True
## False 2638   6
##  True   13 422
```

```
rpart.stats <- get_stats(confMat)
rpart.stats
```

```
##      name      value
## 1 accuracy 0.993829165
## 2 error rate 0.006170835
## 3 precision 0.985981308
## 4 sensitivity 0.970114943
## 5 specificity 0.997730711
## 6 F-measure 0.977983778
## 7 Matthew's CC 0.974439117
```

Random Forest Analysis

1. Make sure that the model assumptions, if any, are satisfied.

No model assumptions of random forest to be satisfied?

2. Assess the model fit and perform diagnostics, if appropriate.

Both methods appear to give identical results. Note the identical accuracy and kappa ratings across all cp tuning parameters in the plain rpart method...

```

set.seed(1)

nasa.rf <- randomForest(as.factor(Hazardous) ~ ., data=nasa)
nasa.rf

##
## Call:
## randomForest(formula = as.factor(Hazardous) ~ ., data = nasa)
##              Type of random forest: classification
##              Number of trees: 500
## No. of variables tried at each split: 4
##
## OOB estimate of error rate: 0.55%
## Confusion matrix:
##      False True class.error
## False  2640    4 0.001512859
## True     13  422 0.029885057

```

```

set.seed(1)

nasa.CVrf <- train(Hazardous ~ ., data=nasa,
                  method="rf",
                  trControl=trainControl(method="cv", number=10,
                                          savePredictions=TRUE,
                                          classProbs=TRUE))
nasa.CVrf

```

```

## Random Forest
##
## 3079 samples
## 20 predictor
## 2 classes: 'False', 'True'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 2771, 2771, 2772, 2770, 2771, 2772, ...
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa
## 2 0.9948020 0.9782101
## 11 0.9961018 0.9836952
## 20 0.9970769 0.9877668
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 20.

```

3. Identify tuning parameters to be used, if appropriate.

The final model selects an mtry tuning parameter value of 20 in order to maximize accuracy.

4. Identify and interpret the effect of selected variables.

```
varImp(nasa.CVrf)
```


```

## rf variable importance
##
## Overall

```

```
## Minimum.Orbit.Intersection      100.00000
## Absolute.Magnitude              9.35397
## Est.Dia.in.KM.max.              7.86373
## Est.Dia.in.KM.range             7.53785
## Est.Dia.in.KM.min.             6.56980
## Close.Approach.Date             0.49328
## Perihelion.Distance             0.30455
## Aphelion.Dist                   0.25229
## Miss.Dist.in.KM                 0.20076
## Mean.Anomaly                    0.16128
## Eccentricity                    0.14131
## Inclination                     0.12331
## Perihelion.Arg                  0.07658
## Jupiter.Tisserand.Invariant      0.06409
## Relative.Velocity.in.KM.per.sec 0.04976
## Asc.Node.Longitude              0.04545
## Mean.Motion                     0.03574
## Orbital.Period                  0.03012
## Semi.Major.Axis                 0.02982
## Orbit.Uncertainty               0.00000
```

```
plot(varImp(nasa.CVrf))
```




5. Evaluate the cross-validated (CV) predictive performance.

```
head(nasa.CVrf$pred)
```

```
##   pred  obs False  True rowIndex mtry Resample
## 1 False False 0.964 0.036      11    2  Fold01
## 2 False False 0.982 0.018      13    2  Fold01
## 3 False False 0.976 0.024      34    2  Fold01
```

```
## 4 False False 0.996 0.004      38    2 Fold01
## 5 True  True 0.164 0.836      66    2 Fold01
## 6 True  True 0.242 0.758      68    2 Fold01

pihatcv.rf <- nasa.CVrf$pred[nasa.CVrf$pred$mtry == 20,]
predcv.rf <- prediction(pihatcv.rf$True, pihatcv.rf$obs)
perfcv.rf <- performance(predcv.rf, "tpr", "fpr")
plot(perfcv.rf)
abline(a=0, b=1, lty=2)
```



```
aucCV.rf <- performance(predcv.rf, "auc")@y.values
aucCV.rf
```

```
## [[1]]
## [1] 0.9986302
```

```
confMat <- table(pihatcv.rf$obs, pihatcv.rf$pred)
confMat
```

```
##
##      False True
## False 2643   1
## True   8   427
```

```
rf.stats <- get_stats(confMat)
rf.stats
```

```
##      name      value
## 1 accuracy 0.997076973
## 2 error rate 0.002923027
## 3 precision 0.997663551
## 4 sensitivity 0.981609195
## 5 specificity 0.999621785
## 6 F-measure 0.989571263
## 7 Matthew's CC 0.987915632
```