

Asteroid Compiled

Annika Lin

2023-04-04

```
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```
library(mlbench)
library(rattle)
```

```
## Loading required package: tibble
```

```
## Loading required package: bitops
```

```
## Rattle: A free graphical interface for data science with R.
## Version 5.5.1 Copyright (c) 2006-2021 Togaware Pty Ltd.
## Type 'rattle()' to shake, rattle, and roll your data.
```

```
library(randomForest)
```

```
## randomForest 4.7-1.1
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:rattle':
##
##      importance
```

```
## The following object is masked from 'package:ggplot2':
##
##      margin
```

```
library(ROCR)
library(glmnet)
```

```
## Loading required package: Matrix
```

```
##
## Attaching package: 'Matrix'
```

```
## The following object is masked from 'package:bitops':
##
##      %&%
```

```
## Loaded glmnet 4.1-7
```

```
library(car)
```

```
## Loading required package: carData
```

```
library(ResourceSelection)
```

```
## ResourceSelection 0.3-5    2019-07-22
```

```
#import dataset filtered for '2017-04-06'
nasa <- read.csv("~/Documents/Georgetown/Spring23/Statistical Learning & Data Science/Pr
oject/NASA-asteroid-Classification-master/final/nasa.csv")
```

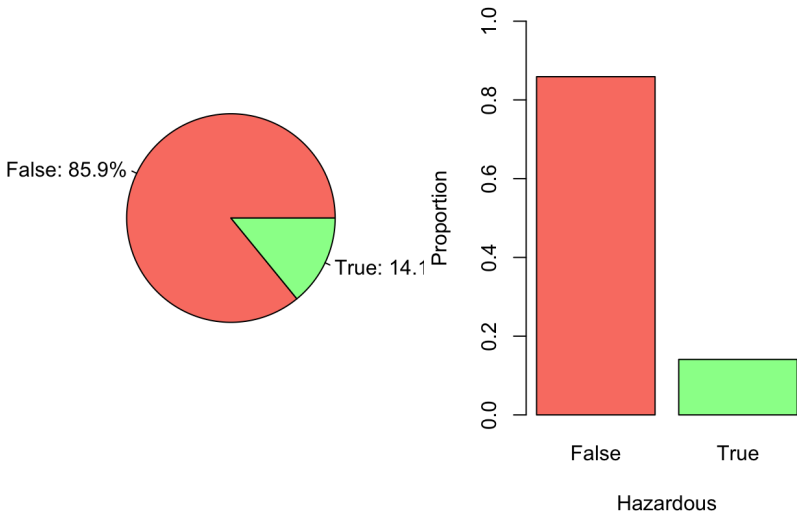
```
nasa <- nasa[ , !(names(nasa) %in% c("X"))]
```

```
#remove outlier
nasa <- nasa[-695,]
```

```
prop.hazardous <- prop.table(table(nasa$Hazardous))
prop.hazardous
```

```
##
##      False      True
## 0.8589994 0.1410006
```

```
par(mfrow=c(1,2))
# pie chart
count.hazardous <- table(nasa$Hazardous)
lbls <- paste(levels(as.factor(nasa$Hazardous)), ":", 
              round(prop.hazardous,3)*100, "%", sep=" ")
pie(count.hazardous, labels=lbls, col=c("salmon", "palegreen"))
# bar plot
barplot(prop.hazardous, xlab="Hazardous", ylab="Proportion", ylim=c(0, 1.0),
        col=c("salmon", "palegreen"))
```



Analyses of Predictors

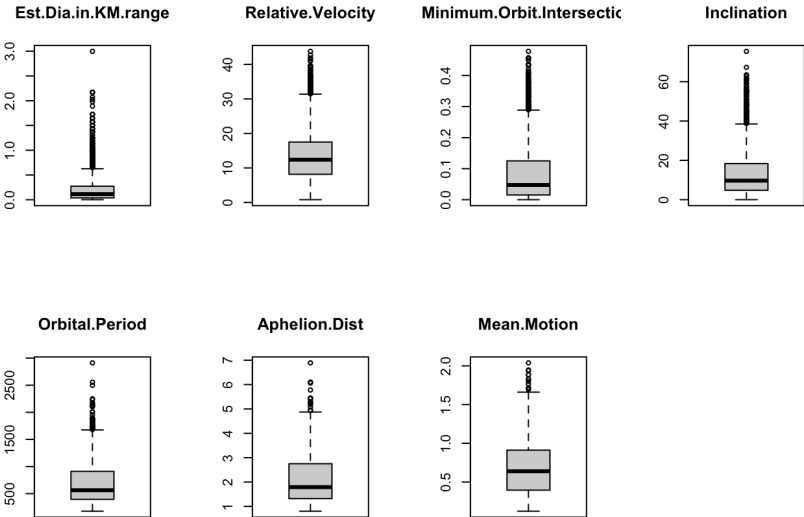
(in search of outliers that might skew analysis)

```
# numerical summary + box plots
summary(nasa)
```

```
## Absolute.Magnitude Est.Dia.in.KM.min. Est.Dia.in.KM.max. Est.Dia.in.KM.range
## Min. :15.20 Min. :0.00101 Min. :0.00226 Min. :0.001249
## 1st Qu.:20.40 1st Qu.:0.03052 1st Qu.:0.06824 1st Qu.:0.037722
## Median :22.30 Median :0.09216 Median :0.20608 Median :0.113919
## Mean :22.53 Mean :0.17185 Mean :0.38428 Mean :0.212423
## 3rd Qu.:24.70 3rd Qu.:0.22108 3rd Qu.:0.49436 3rd Qu.:0.273273
## Max. :32.10 Max. :2.42412 Max. :5.42051 Max. :2.996383
## Close.Approach.Date Relative.Velocity.in.KM.per.sec Miss.Dist.in.KM
## Min. :19950101 Min. : 0.8002 Min. : 26610
## 1st Qu.:20010808 1st Qu.: 8.1680 1st Qu.:16215102
## Median :20070912 Median :12.3870 Median :37033462
## Mean :20066259 Mean :13.6128 Mean :36473114
## 3rd Qu.:20120922 3rd Qu.:17.5079 3rd Qu.:56290664
## Max. :20160908 Max. :43.7899 Max. :74781600
## Orbit.Uncertainty Minimum.Orbit.Intersection Jupiter.Tisserand.Invariant
## Min. :0.0 Min. :0.0000021 Min. :2.196
## 1st Qu.:1.0 1st Qu.:0.0151341 1st Qu.:3.807
## Median :5.0 Median :0.0473452 Median :4.798
## Mean :4.1 Mean :0.0823254 Mean :4.852
## 3rd Qu.:7.0 3rd Qu.:0.1249268 3rd Qu.:5.774
## Max. :9.0 Max. :0.4778910 Max. :9.025
## Eccentricity Semi.Major.Axis Inclination Asc.Node.Longitude
## Min. :0.01296 Min. :0.6159 Min. : 0.01451 Min. : 0.0019
## 1st Qu.:0.24907 1st Qu.:1.0530 1st Qu.: 4.78566 1st Qu.: 83.6648
## Median :0.38236 Median :1.3347 Median : 9.68715 Median :173.5898
## Mean :0.39321 Mean :1.4848 Mean :12.84105 Mean :173.5232
## 3rd Qu.:0.52589 3rd Qu.:1.8386 3rd Qu.:18.38238 3rd Qu.:258.4855
## Max. :0.96026 Max. :3.9908 Max. :75.40667 Max. :359.9059
## Orbital.Period Perihelion.Distance Perihelion.Arg Aphelion.Dist
## Min. : 176.6 Min. :0.08074 Min. : 0.0069 Min. :0.8038
## 1st Qu.: 394.7 1st Qu.:0.67807 1st Qu.: 95.6381 1st Qu.:1.3191
## Median : 563.2 Median :0.87434 Median :188.4906 Median :1.7911
## Mean : 692.8 Mean :0.84323 Mean :184.1273 Mean :2.1264
## 3rd Qu.: 910.6 3rd Qu.:1.01936 3rd Qu.:272.5434 3rd Qu.:2.7517
## Max. :2912.0 Max. :1.29983 Max. :359.9931 Max. :6.8918
## Mean.Anomaly Mean.Motion Hazardous
## Min. : 0.0032 Min. :0.1236 Length:3078
## 1st Qu.: 83.3164 1st Qu.:0.3953 Class :character
## Median :183.8903 Median :0.6392 Mode :character
## Mean :180.1357 Mean :0.6811
## 3rd Qu.:276.3132 3rd Qu.:0.9121
## Max. :359.9180 Max. :2.0390
```

```
par(mfrow=c(2,4))
boxplot(nasa$Est.Dia.in.KM.range, main="Est.Dia.in.KM.range")
boxplot(nasa$Relative.Velocity, main="Relative.Velocity")
boxplot(nasa$Minimum.Orbit.Intersection, main="Minimum.Orbit.Intersection")
boxplot(nasa$Inclination, main="Inclination")
boxplot(nasa$Orbital.Period, main="Orbital.Period")
boxplot(nasa$Aphelion.Dist, main="Aphelion.Dist")
boxplot(nasa$Mean.Motion, main="Mean.Motion")
# to find row number of outlier observations
nasa[which.max(nasa$Est.Dia.in.KM.range),]
```

##	Absolute.Magnitude	Est.Dia.in.KM.min.	Est.Dia.in.KM.max.	
## 2729	15.2	2.424125	5.420508	
##	Est.Dia.in.KM.range	Close.Approach.Date	Relative.Velocity.in.KM.per.sec	
## 2729	2.996383	20141222	23.5172	
##	Miss.Dist.in.KM	Orbit.Uncertainty	Minimum.Orbit.Intersection	
## 2729	45467472	0	0.153116	
##	Jupiter.Tisserand.Invariant	Eccentricity	Semi.Major.Axis	Inclination
## 2729	4.864	0.3462176	1.261473	36.90474
##	Asc.Node.Longitude	Orbital.Period	Perihelion.Distance	Perihelion.Arg
## 2729	111.2844	517.5058	0.8247287	349.1364
##	Aphelion.Dist	Mean.Anomaly	Mean.Motion	Hazardous
## 2729	1.698217	328.5237	0.6956443	False



Predictive Performance Stats

```
get_stats <- function(CM) {
  TP <- CM[2,2]
  FP <- CM[1,2]
  TN <- CM[1,1]
  FN <- CM[2,1]

  acc <- (TP+TN) / (TP+TN+FN+FP)
  err <- (FP+FN) / (TP+TN+FN+FP)
  pre <- (TP) / (TP+FP)
  sen <- (TP) / (TP+FN)
  spe <- (TN) / (TN+FP)
  fme <- (2*pre*sen) / (pre+sen)
  mcc_denom <- sqrt(TP+FP)*sqrt(TP+FN)*sqrt(TN+FP)*sqrt(TN+FN)
  mcc <- (TP*TN - FP*FN) / mcc_denom

  name <- c("accuracy", "error rate", "precision", "sensitivity", "specificity", "F-measure", "Matthew's CC")
  value <- c(acc, err, pre, sen, spe, fme, mcc)
  stats <- data.frame(name, value)

  return (stats)
}
```

Lasso's Penalized Regression

Create the design matrix


```
X = model.matrix(Hazardous ~ ., data=nasa)
Y = as.numeric(nasa$Hazardous=="True")
```

Conduct the cross-validation

```
set.seed(1)
cvfit = cv.glmnet(x=X[, -1], y=Y, family="binomial", type.measure="auc")
cvfit

##
## Call: cv.glmnet(x = X[, -1], y = Y, type.measure = "auc", family = "binomial")
##
## Measure: AUC
##
##      Lambda Index Measure      SE Nonzero
## min 0.001501  48  0.9908 0.0007380      8
## 1se 0.003468  39  0.9902 0.0008217      5

plot(cvfit)
```



Variables selected using λ_{1se}

```
sel.vars <- which(coef(cvfit, s=cvfit$lambda.1se)!=0)[-1]-1
sel.names <- colnames(nasa)[sel.vars]
sel.names
```

```
## [1] "Absolute.Magnitude"      "Est.Dia.in.KM.min."
## [3] "Orbit.Uncertainty"       "Minimum.Orbit.Intersection"
## [5] "Mean.Motion"
```

```
#fit a lasso model using the selected variables
fit.lasso <- glm(as.factor(Hazardous) ~ Absolute.Magnitude + Est.Dia.in.KM.min. + Orbit.
Uncertainty +
               Minimum.Orbit.Intersection + Jupiter.Tisserand.Invariant,
               family="binomial", data=nasa)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(fit.lasso)
```

```
##
## Call:
## glm(formula = as.factor(Hazardous) ~ Absolute.Magnitude + Est.Dia.in.KM.min. +
##      Orbit.Uncertainty + Minimum.Orbit.Intersection + Jupiter.Tisserand.Invariant,
##      family = "binomial", data = nasa)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.21356  -0.02287  -0.00118  -0.00001   2.97254
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      83.95094    6.09896   13.765 < 2e-16 ***
## Absolute.Magnitude    -3.57347    0.26573  -13.448 < 2e-16 ***
## Est.Dia.in.KM.min.   -17.11936    1.54653  -11.070 < 2e-16 ***
## Orbit.Uncertainty     -0.13672    0.04963   -2.755 0.00587 **
## Minimum.Orbit.Intersection -129.67467    8.94106  -14.503 < 2e-16 ***
## Jupiter.Tisserand.Invariant  -0.12295    0.09307   -1.321 0.18645
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 2504.11  on 3077  degrees of freedom
## Residual deviance:  505.29  on 3072  degrees of freedom
## AIC: 517.29
##
## Number of Fisher Scoring iterations: 10
```

```
#address the model assumptions of the lasso model -- multicollinearity
vif(fit.lasso)
```

```
##      Absolute.Magnitude      Est.Dia.in.KM.min.
##      13.402908              8.509262
##      Orbit.Uncertainty  Minimum.Orbit.Intersection
##      1.392265              3.010432
##      Jupiter.Tisserand.Invariant
##      1.159019
```

```
#adjust the model based on multicollinearity issues
#remove Absolute.Magnitude
fit.lasso2 <- glm(as.factor(Hazardous) ~ Est.Dia.in.KM.min. + Orbit.Uncertainty +
               Minimum.Orbit.Intersection + Jupiter.Tisserand.Invariant,
               family="binomial", data=nasa)
summary(fit.lasso2)
```

```
##
## Call:
## glm(formula = as.factor(Hazardous) ~ Est.Dia.in.KM.min. + Orbit.Uncertainty +
##      Minimum.Orbit.Intersection + Jupiter.Tisserand.Invariant,
##      family = "binomial", data = nasa)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.1507  -0.3249  -0.0973  -0.0035   3.1983
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      3.56403    0.41170   8.657 < 2e-16 ***
## Est.Dia.in.KM.min.  4.49382    0.51569   8.714 < 2e-16 ***
## Orbit.Uncertainty  -0.52498    0.03271 -16.049 < 2e-16 ***
## Minimum.Orbit.Intersection -62.34047  4.02315 -15.495 < 2e-16 ***
## Jupiter.Tisserand.Invariant -0.43722    0.06464  -6.764 1.34e-11 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 2504.1  on 3077  degrees of freedom
## Residual deviance: 1180.2  on 3073  degrees of freedom
## AIC: 1190.2
##
## Number of Fisher Scoring iterations: 8
```

```
#assess multicollinearity of the adjusted model
vif(fit.lasso2)
```

```
##           Est.Dia.in.KM.min.      Orbit.Uncertainty
##           2.133114                1.508538
## Minimum.Orbit.Intersection Jupiter.Tisserand.Invariant
##           2.116395                1.110887
```

```
#test the goodness of fit of the adjusted model
hoslem.test(fit.lasso2$y, fit.lasso2$fitted.values)
```

```
##
## Hosmer and Lemeshow goodness of fit (GOF) test
##
## data: fit.lasso2$y, fit.lasso2$fitted.values
## X-squared = 7.974, df = 8, p-value = 0.436
```

```
#adjust the model based on multicollinearity issues
#remove Est.Dia.in.KM.min
fit.lasso3 <- glm(as.factor(Hazardous) ~ Absolute.Magnitude + Orbit.Uncertainty +
      Minimum.Orbit.Intersection + Jupiter.Tisserand.Invariant,
      family="binomial", data=nasa)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(fit.lasso3)
```

```
##
## Call:
## glm(formula = as.factor(Hazardous) ~ Absolute.Magnitude + Orbit.Uncertainty +
##      Minimum.Orbit.Intersection + Jupiter.Tisserand.Invariant,
##      family = "binomial", data = nasa)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.83507  -0.10414  -0.01809  -0.00008   2.57268
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      35.69555    2.16622  16.478 < 2e-16 ***
## Absolute.Magnitude  -1.48606    0.09551 -15.558 < 2e-16 ***
## Orbit.Uncertainty  -0.16696    0.04104  -4.069 4.73e-05 ***
## Minimum.Orbit.Intersection -109.62272  6.86539 -15.967 < 2e-16 ***
## Jupiter.Tisserand.Invariant  -0.13840    0.07946  -1.742  0.0816 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 2504.11  on 3077  degrees of freedom
## Residual deviance:  690.72  on 3073  degrees of freedom
## AIC: 700.72
##
## Number of Fisher Scoring iterations: 9
```

```
#assess multicollinearity of the adjusted model
vif(fit.lasso3)
```

```
##           Absolute.Magnitude      Orbit.Uncertainty
##           3.081047                1.432653
## Minimum.Orbit.Intersection Jupiter.Tisserand.Invariant
##           2.856256                1.123920
```

```
#test the goodness of fit of the adjusted model
hoslem.test(fit.lasso3$y, fit.lasso3$fitted.values)
```

```
##
## Hosmer and Lemeshow goodness of fit (GOF) test
##
## data: fit.lasso3$y, fit.lasso3$fitted.values
## X-squared = 7.0759, df = 8, p-value = 0.5285
```

Fit.lasso3 is the best fit of the lasso model

```
#create the cross-validated model using the selected variables
```

```
set.seed(1)
fit.cv <- train(as.factor(Hazardous) ~ Absolute.Magnitude + Est.Dia.in.KM.min. + Orbit.U
ncertainty +
               Minimum.Orbit.Intersection + Jupiter.Tisserand.Invariant ,
               method = "glm", family = "binomial", trControl = trainControl(method
="cv", number=5,
                                   savePredictions = TRUE),data=nasa)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
fit.cv
```

```
## Generalized Linear Model
##
## 3078 samples
## 5 predictor
## 2 classes: 'False', 'True'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 2462, 2463, 2463, 2462, 2462
## Resampling results:
##
## Accuracy Kappa
## 0.9606937 0.8371297
```

```
#determine the final model
summary(fit.cv$finalModel)
```

```
##
## Call:
## NULL
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.21356  -0.02287  -0.00118  -0.00001   2.97254
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      83.95094    6.09896  13.765 < 2e-16 ***
## Absolute.Magnitude    -3.57347    0.26573 -13.448 < 2e-16 ***
## Est.Dia.in.KM.min.   -17.11936    1.54653 -11.070 < 2e-16 ***
## Orbit.Uncertainty    -0.13672    0.04963  -2.755 0.00587 **
## Minimum.Orbit.Intersection -129.67467    8.94106 -14.503 < 2e-16 ***
## Jupiter.Tisserand.Invariant  -0.12295    0.09307  -1.321 0.18645
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 2504.11 on 3077 degrees of freedom
## Residual deviance: 505.29 on 3072 degrees of freedom
## AIC: 517.29
##
## Number of Fisher Scoring iterations: 10
```

```
#asses the mullitcollinearity of the final model
vif(fit.cv$finalModel)
```

```
##      Absolute.Magnitude      Est.Dia.in.KM.min.
##      13.402908              8.509262
##      Orbit.Uncertainty  Minimum.Orbit.Intersection
##      1.392265              3.010432
##      Jupiter.Tisserand.Invariant
##      1.159019
```

```
#create a new cross-validated model removing Est.Dia.in.KM.min.
fit.cv2 <- train(as.factor(Hazardous) ~ Absolute.Magnitude + Orbit.Uncertainty +
               Minimum.Orbit.Intersection + Jupiter.Tisserand.Invariant ,
               method = "glm", family = "binomial", trControl = trainControl(method
="cv", number=5,
                                   classPr
                                   savePredictions = TRUE),data=nasa)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
fit.cv2
```

```
## Generalized Linear Model
##
## 3078 samples
## 4 predictor
## 2 classes: 'False', 'True'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 2462, 2463, 2463, 2462, 2462
## Resampling results:
##
## Accuracy Kappa
## 0.9496442 0.7864174
```

```
#determine the final model
summary(fit.cv2$finalModel)
```

```
##
## Call:
## NULL
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.83507  -0.10414  -0.01809  -0.00008   2.57268
##
## Coefficients:
##                                Estimate Std. Error z value Pr(>|z|)
## (Intercept)                   35.69555     2.16622  16.478  < 2e-16 ***
## Absolute.Magnitude             -1.48606     0.09551 -15.558  < 2e-16 ***
## Orbit.Uncertainty              -0.16696     0.04104  -4.069 4.73e-05 ***
## Minimum.Orbit.Intersection    -109.62272     6.86539 -15.967  < 2e-16 ***
## Jupiter.Tisserand.Invariant    -0.13840     0.07946  -1.742  0.0816 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 2504.11 on 3077 degrees of freedom
## Residual deviance: 690.72 on 3073 degrees of freedom
## AIC: 700.72
##
## Number of Fisher Scoring iterations: 9
```

```
#assess the multicollinearity of the adjusted cross-validated model
vif(fit.cv2$finalModel)
```

```
##           Absolute.Magnitude           Orbit.Uncertainty
##                3.081047                1.432653
## Minimum.Orbit.Intersection Jupiter.Tisserand.Invariant
##                2.856256                1.123920
```


```
#goodness of fit
hoslem.test(fit.cv2$finalModel$y, fit.cv2$finalModel$fitted.values)
```

```
##
## Hosmer and Lemeshow goodness of fit (GOF) test
##
## data: fit.cv2$finalModel$y, fit.cv2$finalModel$fitted.values
## X-squared = 7.0759, df = 8, p-value = 0.5285
```

```
#assess the predictive performance using the optimized model
pihat <- predict(fit.cv2, type="prob")
head(cbind(nasa$Hazardous, pihat, predict(fit.cv2)))
```

```
## nasa$Hazardous      False      True predict(fit.cv2)
## 1      True 0.6572134 3.427866e-01      False
## 2      False 1.0000000 2.043709e-08      False
## 3      True 0.4556698 5.443302e-01      True
## 4      False 0.9993470 6.530091e-04      False
## 5      True 0.7499228 2.500772e-01      False
## 6      False 0.9871813 1.281874e-02      False
```

```
pred <- prediction(pihat[,2], nasa$Hazardous)
perf <- performance(pred, "tpr", "fpr")
plot(perf)
```



```
auc <- performance(pred, "auc")@y.values
auc
```


```
## [[1]]
## [1] 0.9840792
```

```
#assess the predictive performance using the predictive model
# pihatcv <- fit.cv2$pred
# head(cbind(nasa$Hazardous[pihatcv$rowIndex], pihatcv))
# predcv <- prediction(pihatcv$True, pihatcv$obs)
# perfcv <- performance(predcv, "tpr", "fpr")
# plot(perfcv)
# auccv <- performance(predcv, "auc")@y.values
# auccv
```

```
predprob.lasso <- fit.cv2$pred
head(predprob.lasso)
```

```
##      pred  obs      False      True rowIndex parameter Resample
## 1  True  True 0.44764628 0.552353716      3      none      Fold1
## 2  True  True 0.27155987 0.728440134      7      none      Fold1
## 3 False False 0.99548421 0.004515794     16      none      Fold1
## 4 False False 0.99883589 0.001164109     17      none      Fold1
## 5 False False 0.99724971 0.002750294     20      none      Fold1
## 6  True False 0.04122548 0.958774516     32      none      Fold1
```

```
pred.lasso = prediction(predprob.lasso$True, predprob.lasso$obs)
perf.lasso = performance(pred.lasso, "tpr", "fpr")
plot(perf.lasso)
```




```
auc.lasso = performance(pred.lasso, "auc")@y.values
auc.lasso
```

```
## [[1]]
## [1] 0.983586
```

```
confusionMatrix(predprob.lasso$pred, predprob.lasso$obs, positive="True")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction False True
##           False  2581  92
##           True   63  342
##
##           Accuracy : 0.9496
##           95% CI : (0.9413, 0.9571)
##           No Information Rate : 0.859
##           P-Value [Acc > NIR] : < 2e-16
##
##           Kappa : 0.7861
##
## Mcnemar's Test P-Value : 0.02451
##
##           Sensitivity : 0.7880
##           Specificity : 0.9762
##           Pos Pred Value : 0.8444
##           Neg Pred Value : 0.9656
##           Prevalence : 0.1410
##           Detection Rate : 0.1111
##           Detection Prevalence : 0.1316
##           Balanced Accuracy : 0.8821
##
##           'Positive' Class : True
##
```

```
# Confusion matrix
conf.lasso <- table(predprob.lasso$pred, predprob.lasso$obs)
conf.lasso
```

```
##
##           False True
## False  2581  92
## True   63  342
```

```
lasso.stats <- get_stats(conf.lasso)
lasso.stats
```

```
##           name      value
## 1      accuracy 0.94964263
## 2      error rate 0.05035737
## 3      precision 0.78801843
## 4      sensitivity 0.84444444
## 5      specificity 0.96558174
## 6      F-measure 0.81525626
## 7 Matthew's CC 0.78677483
```

k-Nearest Neighbor

1. Make sure that the model assumptions, if any, are satisfied.

KNN makes no assumptions and has relatively few parameters to specify (k and a distance measure).

2. Assess the model fit and perform diagnostics, if appropriate.

```
# function to normalize data
normalize <- function(x) {
  return ((x - min(x)) / (max(x) - min(x))) }

arr.norm <- apply(nasa[,-21], 2, normalize)

arr.norm <- data.frame(arr.norm, nasa$Hazardous)

colnames(arr.norm)[colnames(arr.norm) == "nasa.Hazardous"] = "Hazardous"
```

```
# 5-fold CV to choose k

set.seed(1)

arr.norm$Hazardous <- as.factor(arr.norm$Hazardous)

fit.knn <- train(Hazardous ~ .,
  method = "knn",
  tuneGrid = expand.grid(k = 1:21),
  trControl = trainControl(method="cv", number=5, savePredictions = TRUE, classProbs = T
RUE),
  metric = "Accuracy",
  data = arr.norm)

fit.knn
```

```
## k-Nearest Neighbors
##
## 3078 samples
## 20 predictor
## 2 classes: 'False', 'True'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 2462, 2463, 2463, 2462, 2462
## Resampling results across tuning parameters:
##
## k Accuracy Kappa
## 1 0.8645222 0.4227282
## 2 0.8534785 0.3734312
## 3 0.8804429 0.4254110
## 4 0.8749203 0.3926299
## 5 0.8814180 0.4071912
## 6 0.8778418 0.3926183
## 7 0.8801167 0.3847148
## 8 0.8797931 0.3796854
## 9 0.8775209 0.3589633
## 10 0.8804466 0.3685311
## 11 0.8801193 0.3552413
## 12 0.8804440 0.3575683
## 13 0.8797936 0.3410263
## 14 0.8797936 0.3371958
## 15 0.8801177 0.3317969
## 16 0.8762200 0.3033070
## 17 0.8788180 0.3181018
## 18 0.8781718 0.3167519
## 19 0.8794684 0.3145760
## 20 0.8768699 0.2957819
## 21 0.8791442 0.3117543
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was k = 5.
```


5 neighbors are used in the final model.

10 most important variables are:

```
imp.knn <- rownames(varImp(fit.knn)$importance)
sel.knn <- imp.knn[order(varImp(fit.knn)$importance[,1], decreasing=T)][1:10]
sel.knn
```

```
## [1] "Absolute.Magnitude" "Est.Dia.in.KM.min."
## [3] "Est.Dia.in.KM.max." "Est.Dia.in.KM.range"
## [5] "Orbit.Uncertainty" "Minimum.Orbit.Intersection"
## [7] "Relative.Velocity.in.KM.per.sec" "Perihelion.Distance"
## [9] "Eccentricity" "Close.Approach.Date"
```

```
pihatcv.knn <- fit.knn$pred[fit.knn$pred$k == 5,]
pred <- prediction(pihatcv.knn$True, pihatcv.knn$obs)
perf <- performance(pred, "tpr", "fpr")
plot(perf)
```




```
# Area under ROC curve (AUC) = concordance index
auc.perf = performance(pred, "auc")
knn_auc <- auc.perf@y.values
knn_auc
```


```
## [[1]]
## [1] 0.8553424
```

KNN has an AUC of 0.8553424.

```
plot(fit.knn$results[,1], 1-fit.knn$results[,2], type="l",
     xlab="k, number of neighbors", ylab="error rate")
```



```
plot(varImp(fit.knn), main="kNN variable importance")
```



```
varImp(fit.knn)
```

```
## ROC curve variable importance
##
##
## Est.Dia.in.KM.range      Importance
## Est.Dia.in.KM.range      100.0000
## Est.Dia.in.KM.max.      100.0000
## Est.Dia.in.KM.min.      100.0000
## Absolute.Magnitude      100.0000
## Orbit.Uncertainty        93.7065
## Minimum.Orbit.Intersection 76.2993
## Relative.Velocity.in.KM.per.sec 58.8524
## Perihelion.Distance      58.3352
## Eccentricity             53.1304
## Close.Approach.Date      26.3045
## Aphelion.Dist            21.3395
## Jupiter.Tisserand.Invariant 12.2465
## Miss.Dist.in.KM          10.0205
## Mean.Anomaly             10.0178
## Mean.Motion              7.0710
## Semi.Major.Axis          7.0710
## Orbital.Period           7.0710
## Inclination              3.4734
## Perihelion.Arg           0.9156
## Asc.Node.Longitude       0.0000
```

```
predprob.knn <- fit.knn$pred
predprob.knn <- predprob.knn[predprob.knn$k==as.numeric(fit.knn$bestTune),]
```

```
pred.knn = prediction(predprob.knn$True, predprob.knn$obs)
perf.knn = performance(pred.knn, "tpr", "fpr")
knn.cut <- cbind(unlist(perf.knn@y.values), unlist(perf.knn@x.values), unlist(perf.knn@alpha.values))
knn.cut[knn.cut[,1]>0.7 & knn.cut[,2]<0.3,]
```

```
## [1] 0.8686636 0.2443268 0.2000000
```

```
knn.cut[knn.cut[,3] == 0.5,]
```

```
##      [,1] [,2] [,3]
```

```
confusionMatrix(predprob.knn$pred, predprob.knn$obs, positive="True")
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction False True
##      False  2551  272
##      True   93   162
##
##              Accuracy : 0.8814
##              95% CI : (0.8695, 0.8926)
##      No Information Rate : 0.859
##      P-Value [Acc > NIR] : 0.0001443
##
##              Kappa : 0.4085
##
##      Mcnemar's Test P-Value : < 2.2e-16
##
##              Sensitivity : 0.37327
##              Specificity : 0.96483
##      Pos Pred Value : 0.63529
##      Neg Pred Value : 0.90365
##              Prevalence : 0.14100
##      Detection Rate : 0.05263
##      Detection Prevalence : 0.08285
##      Balanced Accuracy : 0.66905
##
##      'Positive' Class : True
##
```

```
# Confusion matrix
conf.knn <- table(predprob.knn$pred, predprob.knn$obs)
conf.knn
```

```
##
##              False True
##      False  2551  272
##      True   93   162
```

```
knn.stats <- get_stats(conf.knn)
knn.stats
```

```
##      name      value
## 1 accuracy 0.8814165
## 2 error rate 0.1185835
## 3 precision 0.3732719
## 4 sensitivity 0.6352941
## 5 specificity 0.9036486
## 6 F-measure 0.4702467
## 7 Matthew's CC 0.4268670
```

Classification Tree Analysis

1. Make sure that the model assumptions, if any, are satisfied.

No model assumptions of decision trees to be satisfied? We are fitting a classification tree as opposed to a regression tree because the response variable is categorical and binary.

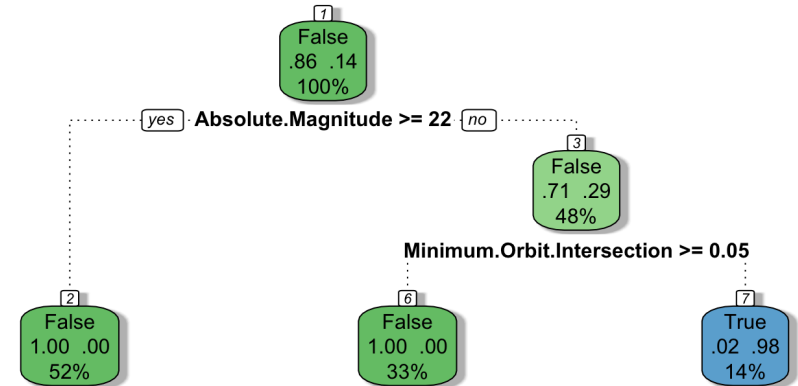
2. Assess the model fit and perform diagnostics, if appropriate.

Both methods appear to give identical results. Note the identical accuracy and kappa ratings across all cp tuning parameters in the plain rpart method...

```
set.seed(1)
# rpart
nasa.CVrpart <- train(Hazardous ~ ., data=nasa,
                      method="rpart",
                      tuneGrid = expand.grid(cp=seq(0.005, 0.05, length=10)),
                      trControl=trainControl(method="cv", number=5,
                                              savePredictions=TRUE,
                                              classProbs=TRUE,
                                              selectionFunction = "oneSE"))
nasa.CVrpart
```

```
## CART
##
## 3078 samples
## 20 predictor
## 2 classes: 'False', 'True'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 2462, 2463, 2463, 2462, 2462
## Resampling results across tuning parameters:
##
##  cp      Accuracy   Kappa
##  0.005  0.9948026  0.978491
##  0.010  0.9948026  0.978543
##  0.015  0.9948026  0.978543
##  0.020  0.9948026  0.978543
##  0.025  0.9948026  0.978543
##  0.030  0.9948026  0.978543
##  0.035  0.9948026  0.978543
##  0.040  0.9948026  0.978543
##  0.045  0.9948026  0.978543
##  0.050  0.9948026  0.978543
##
## Accuracy was used to select the optimal model using the one SE rule.
## The final value used for the model was cp = 0.05.
```

```
# print tree
fancyRpartPlot(nasa.CVrpart$finalModel)
```



Rattle 2023-Apr-29 14:42:37 annikalin

3. Identify tuning parameters to be used, if appropriate.


If the plain rpart method is utilized, it selects a cp (complexity parameter) value of 0.05 to maximize accuracy for the final model. Note that the accuracy and kappa values are identical for each of the cp values tried by the model, so any choice within that range should produce comparable results.

4. Identify and interpret the effect of selected variables.

```
# variable importance (rpart)
varImp(nasa.CVrpart)
```

```
## rpart variable importance
##
##
## Minimum.Orbit.Intersection      Overall
##                               100.000
## Est.Dia.in.KM.min.             17.973
## Absolute.Magnitude             17.973
## Est.Dia.in.KM.max.             17.973
## Est.Dia.in.KM.range            17.973
## Perihelion.Distance            12.601
## Miss.Dist.in.KM                8.932
## Inclination                    6.388
## Relative.Velocity.in.KM.per.sec 2.229
## Jupiter.Tisserand.Invariant     0.000
## Orbital.Period                 0.000
## Mean.Motion                   0.000
## Mean.Anomaly                   0.000
## Perihelion.Arg                 0.000
## Close.Approach.Date            0.000
## Semi.Major.Axis                0.000
## Asc.Node.Longitude             0.000
## Eccentricity                   0.000
## Aphelion.Dist                  0.000
## Orbit.Uncertainty              0.000
```

```
plot(varImp(nasa.CVrpart))
```




5. Evaluate the cross-validated (CV) predictive performance.

```
head(nasa.CVrpart$pred)
```

```
##   pred  obs rowIndex    False      True    cp Resample
## 1  True  True      5  0.01436782 0.985632184 0.005   Fold1
## 2 False False     11  1.00000000 0.000000000 0.005   Fold1
## 3 False False     12  1.00000000 0.000000000 0.005   Fold1
## 4 False False     13  0.99686766 0.003132341 0.005   Fold1
## 5 False False     25  0.99686766 0.003132341 0.005   Fold1
## 6 False False     30  0.99686766 0.003132341 0.005   Fold1
```

```
pihatcv.rpart <- nasa.CVrpart$pred[nasa.CVrpart$pred$cp == 0.05,]
predcv.rpart <- prediction(pihatcv.rpart$True, pihatcv.rpart$obs)
perfcv.rpart <- performance(predcv.rpart, "tpr", "fpr")
plot(perfcv.rpart)
abline(a=0, b=1, lty=2)
```



```
aucCV.rpart <- performance(predcv.rpart, "auc")@y.values
aucCV.rpart
```

```
## [[1]]
## [1] 0.9917045
```

```
confMat <- table(pihatcv.rpart$obs, pihatcv.rpart$pred)
confMat
```

```
##
##           False True
## False  2637     7
## True    9    425
```

```
rpart.stats <- get_stats(confMat)
rpart.stats
```

```
##           name      value
## 1    accuracy 0.994801819
## 2   error rate 0.005198181
## 3    precision 0.983796296
## 4   sensitivity 0.979262673
## 5    specificity 0.997352496
## 6      F-measure 0.981524249
## 7 Matthew's CC 0.978503227
```

Random Forest Analysis

1. Make sure that the model assumptions, if any, are satisfied.

No model assumptions of random forest to be satisfied?

2. Assess the model fit and perform diagnostics, if appropriate.

```
set.seed(1)
nasa.rf <- randomForest(as.factor(Hazardous) ~ ., data=nasa)
nasa.rf
```

```
##
## Call:
## randomForest(formula = as.factor(Hazardous) ~ ., data = nasa)
##           Type of random forest: classification
##           Number of trees: 500
## No. of variables tried at each split: 4
##
## OOB estimate of error rate: 0.49%
## Confusion matrix:
##           False True class.error
## False  2640     4 0.001512859
## True    11   423 0.025345622
```

```
set.seed(1)
nasa.CVrf <- train(Hazardous ~ ., data=nasa,
                  method="rf",
                  trControl=trainControl(method="cv", number=5,
                                         savePredictions=TRUE,
                                         classProbs=TRUE))
nasa.CVrf
```

```
## Random Forest
##
## 3078 samples
## 20 predictor
## 2 classes: 'False', 'True'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 2462, 2463, 2463, 2462, 2462
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa
## 2 0.9944774 0.9769498
## 11 0.9954519 0.9811428
## 20 0.9961013 0.9838326
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 20.
```

3. Identify tuning parameters to be used, if appropriate.


The final model selects an mtry tuning parameter value of 20 in order to maximize accuracy.

4. Identify and interpret the effect of selected variables.

```
varImp(nasa.CVrf)
```

```
## rf variable importance
##
## Overall
## Minimum.Orbit.Intersection 100.00000
## Est.Dia.in.KM.min. 8.31765
## Est.Dia.in.KM.max. 7.54595
## Est.Dia.in.KM.range 7.19967
## Absolute.Magnitude 6.50834
## Close.Approach.Date 0.50215
## Perihelion.Distance 0.31000
## Aphelion.Dist 0.24640
## Miss.Dist.in.KM 0.14536
## Mean.Anomaly 0.14268
## Eccentricity 0.13382
## Inclination 0.09681
## Perihelion.Arg 0.07256
## Jupiter.Tisserand.Invariant 0.04657
## Orbital.Period 0.03027
## Relative.Velocity.in.KM.per.sec 0.02670
## Mean.Motion 0.01485
## Asc.Node.Longitude 0.01116
## Orbit.Uncertainty 0.01089
## Semi.Major.Axis 0.00000
```

```
plot(varImp(nasa.CVrf))
```




5. Evaluate the cross-validated (CV) predictive performance.

```
head(nasa.CVrf$pred)
```

```
## pred obs False True rowIndex mtry Resample
## 1 True True 0.200 0.800 5 2 Fold1
## 2 False False 0.970 0.030 11 2 Fold1
## 3 False False 0.770 0.230 12 2 Fold1
## 4 False False 0.980 0.020 13 2 Fold1
## 5 False False 0.976 0.024 25 2 Fold1
## 6 False False 0.998 0.002 30 2 Fold1
```

```
pihatcv.rf <- nasa.CVrf$pred[nasa.CVrf$pred$mtry == 20,]
predcv.rf <- prediction(pihtcv.rf$True, pihatcv.rf$obs)
perfcv.rf <- performance(predcv.rf, "tpr", "fpr")
plot(perfcv.rf)
abline(a=0, b=1, lty=2)
```

```
aucCV.rf <- performance(predcv.rf, "auc")@y.values
aucCV.rf
```

```
## [[1]]
## [1] 0.9986715
```

```
confMat <- table(pihatcv.rf$obs, pihatcv.rf$pred)
confMat
```

```
##
##      False True
## False  2640   4
## True     8  426
```

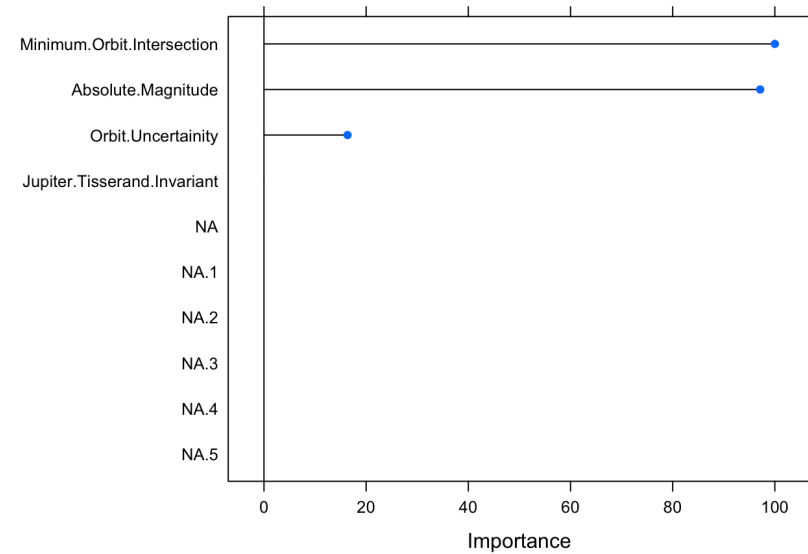
```
rf.stats <- get_stats(confMat)
rf.stats
```

```
##      name      value
## 1  accuracy 0.996101365
## 2  error rate 0.003898635
## 3  precision 0.990697674
## 4  sensitivity 0.981566820
## 5  specificity 0.998487141
## 6  F-measure 0.986111111
## 7  Matthew's CC 0.983857862
```


Model Comparisons

variable importance comparison


```
par(mfrow=c(2,2))
plot(varImp(fit.cv2), top=10)
```




```
plot(varImp(fit.knn), top=10)
```



```
plot(varImp(nasa.CVrpart), top=10)
```



```
plot(varImp(nasa.CVrf), top=10)
```



stats table

lasso.stats		
##	name	value
## 1	accuracy	0.94964263
## 2	error rate	0.05035737
## 3	precision	0.78801843
## 4	sensitivity	0.84444444
## 5	specificity	0.96558174
## 6	F-measure	0.81525626
## 7	Matthew's CC	0.78677483

knn.stats		
-----------	--	--

##	name	value
## 1	accuracy	0.8814165
## 2	error rate	0.1185835
## 3	precision	0.3732719
## 4	sensitivity	0.6352941
## 5	specificity	0.9036486
## 6	F-measure	0.4702467
## 7	Matthew's CC	0.4268670

rpart.stats		
##	name	value
## 1	accuracy	0.994801819
## 2	error rate	0.005198181
## 3	precision	0.983796296
## 4	sensitivity	0.979262673
## 5	specificity	0.997352496
## 6	F-measure	0.981524249
## 7	Matthew's CC	0.978503227

rf.stats		
##	name	value
## 1	accuracy	0.996101365
## 2	error rate	0.003898635
## 3	precision	0.990697674
## 4	sensitivity	0.981566820
## 5	specificity	0.998487141
## 6	F-measure	0.986111111
## 7	Matthew's CC	0.983857862

```
df_merge <- merge(lasso.stats,knn.stats,by="name")
colnames(df_merge)[colnames(df_merge) == "value.x"] = "Lasso"
colnames(df_merge)[colnames(df_merge) == "value.y"] = "kNN"

df_merge <- merge(df_merge,rpart.stats,by="name")
df_merge <- merge(df_merge,rf.stats,by="name")
colnames(df_merge)[colnames(df_merge) == "value.x"] = "Classification Tree"
colnames(df_merge)[colnames(df_merge) == "value.y"] = "Random Forest"

df_merge
```

##	name	Lasso	kNN	Classification Tree	Random Forest
## 1	accuracy	0.94964263	0.8814165	0.994801819	0.996101365
## 2	error rate	0.05035737	0.1185835	0.005198181	0.003898635
## 3	F-measure	0.81525626	0.4702467	0.981524249	0.986111111
## 4	Matthew's CC	0.78677483	0.4268670	0.978503227	0.983857862
## 5	precision	0.78801843	0.3732719	0.983796296	0.990697674
## 6	sensitivity	0.84444444	0.6352941	0.979262673	0.981566820
## 7	specificity	0.96558174	0.9036486	0.997352496	0.998487141

```
imp.knn <- rownames(varImp(fit.knn)$importance)
sel.knn <- imp.knn[order(varImp(fit.knn)$importance[,1], decreasing=T)][1:10]

imp.rp <- rownames(varImp(nasa.CVrpart)$importance)
sel.rp <- imp.rp[order(varImp(nasa.CVrpart)$importance[,1], decreasing=T)][1:10]

imp.rf <- rownames(varImp(nasa.CVrf)$importance)
sel.rf <- imp.rf[order(varImp(nasa.CVrf)$importance[,1], decreasing=T)][1:10]

intersect(sel.names, intersect(sel.knn, intersect(sel.rp, sel.rf)))
```

```
## [1] "Absolute.Magnitude"          "Est.Dia.in.KM.min."
## [3] "Minimum.Orbit.Intersection"
```