

Asteriod Data Analysis

Annika Lin, Hannah Norman, and Madeline Pfister

2023-05-02

Load Packages and Functions

```
library(caret)

## Loading required package: ggplot2
## Loading required package: lattice
library(mlbench)
library(rattle)

## Loading required package: tibble
## Loading required package: bitops
## Rattle: A free graphical interface for data science with R.
## Version 5.5.1 Copyright (c) 2006-2021 Togaware Pty Ltd.
## Type 'rattle()' to shake, rattle, and roll your data.
library(randomForest)

## randomForest 4.7-1.1
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
## The following object is masked from 'package:rattle':
##
##     importance
## The following object is masked from 'package:ggplot2':
##
##     margin
library(ROCR)
library(glmnet)

## Loading required package: Matrix
##
## Attaching package: 'Matrix'
## The following object is masked from 'package:bitops':
##
##     %&%
## Loaded glmnet 4.1-7
```

```
library(car)
```

```
## Loading required package: carData
```

```
library(ResourceSelection)
```

```
## ResourceSelection 0.3-5    2019-07-22
```

Predictive Performance Stats

We created a function in which the output is a table summarizing all of the predictive performance stats. This function was used throughout the analysis.

```
get_stats <- function(CM) {  
  TP <- CM[2,2]  
  FP <- CM[2,1]  
  TN <- CM[1,1]  
  FN <- CM[1,2]  
  
  acc <- (TP+TN) / (TP+TN+FN+FP)  
  err <- (FP+FN) / (TP+TN+FN+FP)  
  pre <- (TP) / (TP+FP)  
  sen <- (TP) / (TP+FN)  
  spe <- (TN) / (TN+FP)  
  fme <- (2*pre*sen) / (pre+sen)  
  mcc_denom <- sqrt(TP+FP)*sqrt(TP+FN)*sqrt(TN+FP)*sqrt(TN+FN)  
  mcc <- (TP*TN - FP*FN) / mcc_denom  
  
  name <- c("accuracy", "error rate", "precision", "sensitivity", "specificity",  
            "F-measure", "Matthew's CC")  
  value <- c(acc, err, pre, sen, spe, fme, mcc)  
  stats <- data.frame(name, value)  
  
  return (stats)  
}
```

Data Cleaning and Initial Analysis

We began by choosing one date to focus on for our data set. The data we chose was April 6, 2017, so we filtered our data to only include observations from that date.

```
#import dataset filtered for '2017-04-06'  
nasa_outlier <- read.csv("nasa_outlier.csv")
```

Analysis of Predictors

We did a simple numerical summary and box plots to determine if there are any large outliers in the data set.

```
# numerical summary + box plots  
summary(nasa_outlier)
```

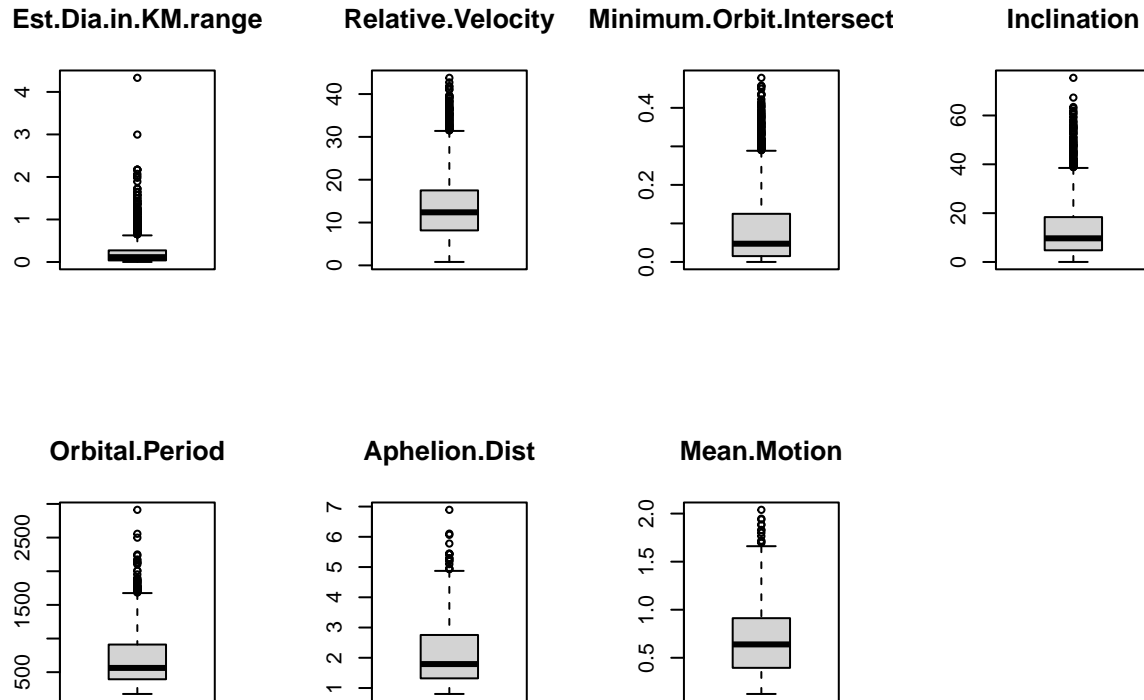
```
## Absolute.Magnitude Est.Dia.in.KM.min. Est.Dia.in.KM.max. Est.Dia.in.KM.range  
## Min. :14.40 Min. :0.001011 Min. :0.00226 Min. :0.001249  
## 1st Qu.:20.40 1st Qu.:0.030518 1st Qu.:0.06824 1st Qu.:0.037722  
## Median :22.30 Median :0.092163 Median :0.20608 Median :0.113919  
## Mean :22.53 Mean :0.172936 Mean :0.38670 Mean :0.213761
```

```
## 3rd Qu.:24.70      3rd Qu.:0.221083      3rd Qu.:0.49436      3rd Qu.:0.273273
## Max. :32.10      Max. :3.503926      Max. :7.83502      Max. :4.331091
## Close.Approach.Date Relative.Velocity.in.KM.per.sec Miss.Dist.in.KM
## Min. :19950101      Min. : 0.8002      Min. : 26610
## 1st Qu.:20010765      1st Qu.: 8.1683      1st Qu.:16224266
## Median :20070908      Median :12.3900      Median :37032388
## Mean :20066238      Mean :13.6152      Mean :36468199
## 3rd Qu.:20120922      3rd Qu.:17.5084      3rd Qu.:56281652
## Max. :20160908      Max. :43.7899      Max. :74781600
## Orbit.Uncertainty Minimum.Orbit.Intersection Jupiter.Tisserand.Invariant
## Min. :0.000      Min. :0.0000021      Min. :2.196
## 1st Qu.:1.000      1st Qu.:0.0151384      1st Qu.:3.804
## Median :5.000      Median :0.0473248      Median :4.798
## Mean :4.099      Mean :0.0823078      Mean :4.852
## 3rd Qu.:7.000      3rd Qu.:0.1248985      3rd Qu.:5.774
## Max. :9.000      Max. :0.4778910      Max. :9.025
## Eccentricity Semi.Major.Axis Inclination Asc.Node.Longitude
## Min. :0.01296      Min. :0.6159      Min. : 0.01451      Min. : 0.0019
## 1st Qu.:0.24918      1st Qu.:1.0530      1st Qu.: 4.78919      1st Qu.: 83.6762
## Median :0.38253      Median :1.3349      Median : 9.68518      Median :173.6808
## Mean :0.39329      Mean :1.4850      Mean :12.83906      Mean :173.5626
## 3rd Qu.:0.52595      3rd Qu.:1.8388      3rd Qu.:18.38036      3rd Qu.:258.6316
## Max. :0.96026      Max. :3.9908      Max. :75.40667      Max. :359.9059
## Orbital.Period Perihelion.Distance Perihelion.Arg Aphelion.Dist
## Min. : 176.6      Min. :0.08074      Min. : 0.0069      Min. :0.8038
## 1st Qu.: 394.7      1st Qu.:0.67808      1st Qu.: 95.6430      1st Qu.:1.3191
## Median : 563.3      Median :0.87390      Median :188.5239      Median :1.7918
## Mean : 692.9      Mean :0.84320      Mean :184.1443      Mean :2.1268
## 3rd Qu.: 910.7      3rd Qu.:1.01935      3rd Qu.:272.5059      3rd Qu.:2.7545
## Max. :2912.0      Max. :1.29983      Max. :359.9931      Max. :6.8918
## Mean.Anomaly Mean.Motion Hazardous
## Min. : 0.0032      Min. :0.1236      Length:3079
## 1st Qu.: 83.5492      1st Qu.:0.3953      Class :character
## Median :183.9847      Median :0.6391      Mode :character
## Mean :180.1871      Mean :0.6810
## 3rd Qu.:276.3719      3rd Qu.:0.9121
## Max. :359.9180      Max. :2.0390
```

```
par(mfrow=c(2,4))
boxplot(nasa_outlier$Est.Dia.in.KM.range, main="Est.Dia.in.KM.range")
boxplot(nasa_outlier$Relative.Velocity, main="Relative.Velocity")
boxplot(nasa_outlier$Minimum.Orbit.Intersection, main="Minimum.Orbit.Intersection")
boxplot(nasa_outlier$Inclination, main="Inclination")
boxplot(nasa_outlier$Orbital.Period, main="Orbital.Period")
boxplot(nasa_outlier$Aphelion.Dist, main="Aphelion.Dist")
boxplot(nasa_outlier$Mean.Motion, main="Mean.Motion")
# to find row number of outlier observations
nasa_outlier[which.max(nasa_outlier$Est.Dia.in.KM.range),]
```

```
## Absolute.Magnitude Est.Dia.in.KM.min. Est.Dia.in.KM.max.
## 695 14.4 3.503926 7.835018
## Est.Dia.in.KM.range Close.Approach.Date Relative.Velocity.in.KM.per.sec
## 695 4.331091 20001222 21.19854
## Miss.Dist.in.KM Orbit.Uncertainty Minimum.Orbit.Intersection
## 695 21340634 0 0.0282524
```

```
## Jupiter.Tisserand.Invariant Eccentricity Semi.Major.Axis Inclination
## 695 3.573 0.6343498 1.982214 6.705068
## Asc.Node.Longitude Orbital.Period Perihelion.Distance Perihelion.Arg
## 695 294.8956 1019.352 0.7247968 236.3403
## Aphelion.Dist Mean.Anomaly Mean.Motion Hazardous
## 695 3.239631 338.28 0.3531656 True
```



In an effort to identify extreme outliers, we selected 7 predictors whose numerical summaries suggested the existence of extreme outliers. Observation 695 was the only outlier were deemed extreme enough to remove.

```
#remove outlier
nasa <- nasa_outlier[-695,]
```

Create a Proportion Table


We wanted to create a proportion table to get an understanding of how many asteroids were classified as hazardous and how many were classified as non-hazardous before beginning our analysis.

```
prop.hazardous <- prop.table(table(nasa$Hazardous))
prop.hazardous
```


```
##
## False True
## 0.8589994 0.1410006
```

We found that the majority of our asteroids (85.9%) were classified as non-hazardous. All of our observations were classified as either hazardous or non-hazardous. This is also shown by the visualizations below.

```
# pie chart
count.hazardous <- table(nasa$Hazardous)
lbls <- paste(levels(as.factor(nasa$Hazardous)), ":", " ",
              round(prop.hazardous,3)*100, "%", sep="")
pie(count.hazardous, labels=lbls, col=c("salmon", "palegreen"))
```



```
# bar plot
barplot(prop.hazardous, xlab="Hazardous", ylab="Proportion", ylim=c(0, 1.0),
        col=c("salmon", "palegreen"))
```



Data Analysis

Lasso's Penalized Regression


Create the design matrix.

```
X = model.matrix(Hazardous ~ ., data=nasa)
Y = as.numeric(nasa$Hazardous=="True")
```

Conduct the cross-validation.

```
set.seed(1)
cvfit = cv.glmnet(x=X[,-1], y=Y, family="binomial", type.measure="auc")
cvfit

##
## Call:  cv.glmnet(x = X[, -1], y = Y, type.measure = "auc", family = "binomial")
##
## Measure: AUC
##
##      Lambda Index Measure      SE Nonzero
## min 0.001501   48  0.9908 0.0007380      8
## 1se 0.003468   39  0.9902 0.0008217      5
plot(cvfit)
```



Determine which variables were selected using λ_{1se} . These are the variables with which the lasso models will be built.

```
sel.vars <- which(coef(cvfit, s=cvfit$lambda.1se)!=0)[-1]-1
sel.names <- colnames(nasa)[sel.vars]
sel.names
```

```
## [1] "Absolute.Magnitude" "Est.Dia.in.KM.min."
```

```
## [3] "Orbit.Uncertainty"          "Minimum.Orbit.Intersection"
## [5] "Mean.Motion"
```

We built our initial lasso model using the variables selected.

```
#fit a lasso model using the selected variables
fit.lasso <- glm(as.factor(Hazardous) ~ Absolute.Magnitude + Est.Dia.in.KM.min. +
                Orbit.Uncertainty + Minimum.Orbit.Intersection +
                Mean.Motion, family="binomial", data=nasa)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(fit.lasso)
```

```
##
## Call:
## glm(formula = as.factor(Hazardous) ~ Absolute.Magnitude + Est.Dia.in.KM.min. +
##      Orbit.Uncertainty + Minimum.Orbit.Intersection + Mean.Motion,
##      family = "binomial", data = nasa)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.22020  -0.02263  -0.00116  -0.00001   2.99174
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      83.83471     6.10198  13.739 < 2e-16 ***
## Absolute.Magnitude      -3.57931     0.26542  -13.486 < 2e-16 ***
## Est.Dia.in.KM.min.     -17.12997     1.54537  -11.085 < 2e-16 ***
## Orbit.Uncertainty       -0.13717     0.04917   -2.790  0.00528 **
## Minimum.Orbit.Intersection -129.95192     8.97244  -14.483 < 2e-16 ***
## Mean.Motion          -0.49633     0.33275   -1.492  0.13580
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 2504.11  on 3077  degrees of freedom
## Residual deviance:  504.81  on 3072  degrees of freedom
## AIC: 516.81
##
## Number of Fisher Scoring iterations: 10
```

The model assumptions for independence, normality, and no influential points are satisfied by the data. However, we also need to check for multicollinearity to ensure that all model assumptions are satisfied.

```
#address the model assumptions of the lasso model -- multicollinearity
vif(fit.lasso)
```

```
##      Absolute.Magnitude      Est.Dia.in.KM.min.
##      13.374033           8.520165
##      Orbit.Uncertainty Minimum.Orbit.Intersection
##      1.365700           3.027934
##      Mean.Motion
##      1.130061
```

Both Absolute.Magnitude and Est.Dia.in.KM.min. have a $VIF > 5$ indicating multicollinearity. We decided to initially adjust our model by removing Absolute.Magnitude as it has a larger VIF.

```
#adjust the model based on multicollinearity issues
#remove Absolute.Magnitude
fit.lasso2 <- glm(as.factor(Hazardous) ~ Est.Dia.in.KM.min. + Orbit.Uncertainty +
                  Minimum.Orbit.Intersection + Mean.Motion,
                  family="binomial", data=nasa)
summary(fit.lasso2)
```

```
##
## Call:
## glm(formula = as.factor(Hazardous) ~ Est.Dia.in.KM.min. + Orbit.Uncertainty +
##      Minimum.Orbit.Intersection + Mean.Motion, family = "binomial",
##      data = nasa)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.1854  -0.3306  -0.0984  -0.0036   3.1953
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      2.34246    0.27245   8.598 < 2e-16 ***
## Est.Dia.in.KM.min.  4.65616    0.51422   9.055 < 2e-16 ***
## Orbit.Uncertainty   -0.51699    0.03237 -15.974 < 2e-16 ***
## Minimum.Orbit.Intersection -62.10827    4.01705 -15.461 < 2e-16 ***
## Mean.Motion        -1.41537    0.23374  -6.055 1.4e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 2504.1  on 3077  degrees of freedom
## Residual deviance: 1189.8  on 3073  degrees of freedom
## AIC: 1199.8
##
## Number of Fisher Scoring iterations: 8
```

```
#assess multicollinearity of the adjusted model
vif(fit.lasso2)
```

```
##      Est.Dia.in.KM.min.      Orbit.Uncertainty
##      2.155630            1.479609
## Minimum.Orbit.Intersection      Mean.Motion
##      2.135797            1.080829
```

Removing Absolute.Magnitude fixed the multicollinearity issues. All assumptions are satisfied for the adjusted model. We then tested how well the model fit the data using the Hosmer and Lemeshow goodness of fit test.

H_0 : the model fits the data well

H_1 : the model does not fit the data well

$\alpha = 0.05$

```
#test the goodness of fit of the adjusted model
hoslem.test(fit.lasso2$y, fit.lasso2$fitted.values)
```

```
##
## Hosmer and Lemeshow goodness of fit (GOF) test
```



```
##
## data: fit.lasso2$y, fit.lasso2$fitted.values
## X-squared = 9.4684, df = 8, p-value = 0.3043
```

There is statistically sufficient evidence ($p = 0.436$, $df = 8$) to conclude that the model fits the data well.

We then replaced Est.Dia.in.KM.min. with Absolute.Magnitude (adjusting for multicollinearity) to see if this model would be a better fit for the data.

```
#adjust the model based on multicollinearity issues
#remove Est.Dia.in.KM.min
fit.lasso3 <- glm(as.factor(Hazardous) ~ Absolute.Magnitude + Orbit.Uncertainty +
                  Minimum.Orbit.Intersection + Mean.Motion,
                  family="binomial", data=nasa)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(fit.lasso3)
```

```
##
## Call:
## glm(formula = as.factor(Hazardous) ~ Absolute.Magnitude + Orbit.Uncertainty +
##      Minimum.Orbit.Intersection + Mean.Motion, family = "binomial",
##      data = nasa)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.84132  -0.10369  -0.01797  -0.00008   2.56744
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      35.50759    2.16373   16.410 < 2e-16 ***
## Absolute.Magnitude      -1.49105    0.09507  -15.684 < 2e-16 ***
## Orbit.Uncertainty       -0.16626    0.04068   -4.087 4.37e-05 ***
## Minimum.Orbit.Intersection -109.77605    6.87597  -15.965 < 2e-16 ***
## Mean.Motion         -0.53944    0.28509   -1.892  0.0585 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 2504.11  on 3077  degrees of freedom
## Residual deviance:  690.15  on 3073  degrees of freedom
## AIC: 700.15
##
## Number of Fisher Scoring iterations: 9
```

```
#assess multicollinearity of the adjusted model
vif(fit.lasso3)
```

```
##      Absolute.Magnitude      Orbit.Uncertainty
##              3.045975              1.406947
## Minimum.Orbit.Intersection      Mean.Motion
##              2.861872              1.094045
```

There are no multicollinearity issues in the adjusted model.

Hosmer and Lemeshow goodness of fit test:

H_0 : the model fits the data well

H_1 : the model does not fit the data well

$$\alpha = 0.05$$

```
#test the goodness of fit of the adjusted model
hoslem.test(fit.lasso3$y, fit.lasso3$fitted.values)
```

```
##
## Hosmer and Lemeshow goodness of fit (GOF) test
##
## data: fit.lasso3$y, fit.lasso3$fitted.values
## X-squared = 7.0768, df = 8, p-value = 0.5284
```

There is statistically sufficient evidence ($p = 0.5285$, $df = 8$) to conclude that the model fits the data well.

Fit.lasso3 is the best fit of the lasso model as the Hosmer and Lemeshow goodness of fit (GOF) test results in a higher p-value for fit.lasso3 than fit.lasso2. This is the model in which we base the cross-validated prediction off of.

```
#create the cross-validated model using the selected variables
set.seed(1)
fit.cv <- train(as.factor(Hazardous) ~ Absolute.Magnitude + Orbit.Uncertainty +
               Minimum.Orbit.Intersection + Mean.Motion ,
               method = "glm", family = "binomial",
               trControl = trainControl(method="cv", number=5,
               savePredictions = TRUE, classProbs = TRUE),data=nasa)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
fit.cv
```

```
## Generalized Linear Model
##
## 3078 samples
## 4 predictor
## 2 classes: 'False', 'True'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 2462, 2463, 2463, 2462, 2462
## Resampling results:
##
## Accuracy Kappa
## 0.9470499 0.7763952
```

```
#determine the final model
summary(fit.cv$finalModel)
```

```
##
## Call:
## NULL
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.84132  -0.10369  -0.01797  -0.00008   2.56744
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      35.50759     2.16373   16.410 < 2e-16 ***
## Absolute.Magnitude -1.49105     0.09507  -15.684 < 2e-16 ***
## Orbit.Uncertainty  -0.16626     0.04068   -4.087 4.37e-05 ***
## Minimum.Orbit.Intersection -109.77605     6.87597  -15.965 < 2e-16 ***
## Mean.Motion        -0.53944     0.28509   -1.892  0.0585 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 2504.11  on 3077  degrees of freedom
## Residual deviance:  690.15  on 3073  degrees of freedom
## AIC: 700.15
##
## Number of Fisher Scoring iterations: 9
```

Verify that there are no multicollinearity issues in the final model and that the model fits the data well.

```
#asses the mulitcollinearity of the final model
vif(fit.cv$finalModel)
```

```
##      Absolute.Magnitude      Orbit.Uncertainty
##      3.045975              1.406947
## Minimum.Orbit.Intersection      Mean.Motion
##      2.861872              1.094045
```

H_0 : the model fits the data well

H_1 : the model does not fit the data well


$\alpha = 0.05$

```
#goodness of fit
hoslem.test(fit.cv$finalModel$y, fit.cv$finalModel$fitted.values)
```

```
##
## Hosmer and Lemeshow goodness of fit (GOF) test
##
## data:  fit.cv$finalModel$y, fit.cv$finalModel$fitted.values
## X-squared = 7.0768, df = 8, p-value = 0.5284
```

We then assessed the predictive performance of the model by plotting the ROC curve and finding the area under the curve.

```
#plot the ROC curve
pihatcv <- fit.cv$pred
predcv <- prediction(pihatcv$True, pihatcv$obs)
perfcv <- performance(predcv, "tpr", "fpr")
plot(perfcv)
```




```
#find the area under the ROC curve  
auccv <- performance(predcv, "auc")@y.values  
auccv
```

```
## [[1]]  
## [1] 0.9834971
```

Determine the importance of each of the predictors is used in the regression.

```
plot(varImp(fit.cv), top=4)
```



Analyze the data analysis by creating a confusion matrix and generating performance statistics.

```
confusionMatrix(pihatcv$pred, pihatcv$obs, positive="True")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction False True
##      False  2574   93
##      True    70  341
##
##              Accuracy : 0.947
##              95% CI : (0.9385, 0.9547)
##      No Information Rate : 0.859
##      P-Value [Acc > NIR] : < 2e-16
##
##              Kappa : 0.7764
##
##  Mcnemar's Test P-Value : 0.08486
##
##              Sensitivity : 0.7857
##              Specificity : 0.9735
##              Pos Pred Value : 0.8297
##              Neg Pred Value : 0.9651
##              Prevalence : 0.1410
##              Detection Rate : 0.1108
##      Detection Prevalence : 0.1335
##              Balanced Accuracy : 0.8796
##
##              'Positive' Class : True
```

```
##
# Confusion matrix
conf.lasso <- table(pihatcv$pred, pihatcv$obs)
conf.lasso
```

```
##
##           False True
## False  2574   93
## True    70   341
```

```
lasso.stats <- get_stats(conf.lasso)
lasso.stats
```

```
##           name      value
## 1    accuracy 0.94704353
## 2   error rate 0.05295647
## 3    precision 0.82968370
## 4  sensitivity 0.78571429
## 5  specificity 0.97352496
## 6    F-measure 0.80710059
## 7 Matthew's CC 0.77682255
```

k-Nearest Neighbor

The kNN makes no assumptions and has one parameter to specify (k). We normalized our data before performing the analysis.

```
# function to normalize data
normalize <- function(x) {
  return ((x - min(x)) / (max(x) - min(x))) }
arr.norm <- apply(nasa[,21], 2, normalize)
arr.norm <- data.frame(arr.norm, nasa$Hazardous)
colnames(arr.norm)[colnames(arr.norm) == "nasa.Hazardous"] = "Hazardous"
```

We used the 5-fold cross validated model to choose k, and we found that 5 neighbors are used in the final model.


```
# 5-fold CV to choose k
set.seed(1)
arr.norm$Hazardous <- as.factor(arr.norm$Hazardous)
fit.knn <- train(Hazardous ~ .,
  method = "knn",
  tuneGrid = expand.grid(k = 1:21),
  trControl = trainControl(method="cv", number=5, savePredictions = TRUE, classProbs = TRUE),
  metric = "Accuracy",
  data = arr.norm)
fit.knn
```

```
## k-Nearest Neighbors
##
## 3078 samples
## 20 predictor
## 2 classes: 'False', 'True'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 2462, 2463, 2463, 2462, 2462
## Resampling results across tuning parameters:
##
##  k  Accuracy  Kappa
##  1  0.8645222  0.4227282
##  2  0.8534785  0.3734312
##  3  0.8804429  0.4254110
##  4  0.8749203  0.3926299
##  5  0.8814180  0.4071912
##  6  0.8778418  0.3926183
##  7  0.8801167  0.3847148
##  8  0.8797931  0.3796854
##  9  0.8775209  0.3589633
## 10  0.8804466  0.3685311
## 11  0.8801193  0.3552413
## 12  0.8804440  0.3575683
## 13  0.8797936  0.3410263
## 14  0.8797936  0.3371958
## 15  0.8801177  0.3317969
## 16  0.8762200  0.3033070
## 17  0.8788180  0.3181018
## 18  0.8781718  0.3167519
```

```
## 19 0.8794684 0.3145760
## 20 0.8768699 0.2957819
## 21 0.8791442 0.3117543
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was k = 5.
```

We then assessed the predictive performance of the model by plotting the ROC curve and finding the area under the curve, plotting the error rate, and plotting the most important variables.

```
#plot the ROC curve
pihatcv.knn <- fit.knn$pred[fit.knn$pred$k == 5,]
pred <- prediction(pihatcv.knn$True, pihatcv.knn$obs)
perf <- performance(pred, "tpr", "fpr")
plot(perf)
```



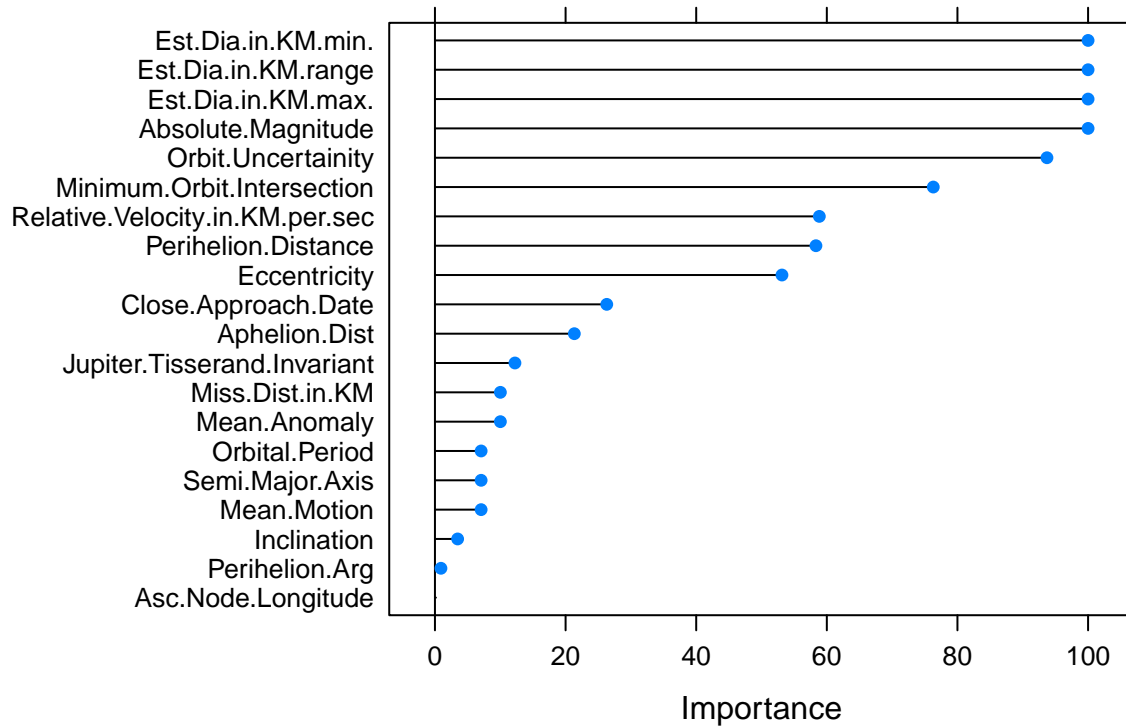
```
# Area under ROC curve (AUC) = concordance index
auc.perf = performance(pred, "auc")
knn_auc <- auc.perf@y.values
knn_auc
```

```
## [[1]]
## [1] 0.8553424
```

The 10 most important variables in the k-Nearest neighbor models we built are:

```
plot(varImp(fit.knn), main="kNN variable importance")
```


kNN variable importance



```
varImp(fit.knn)
```

```
## ROC curve variable importance
##
##                                     Importance
## Est.Dia.in.KM.range                100.0000
## Est.Dia.in.KM.max.                 100.0000
## Est.Dia.in.KM.min.                 100.0000
## Absolute.Magnitude                 100.0000
## Orbit.Uncertainty                   93.7065
## Minimum.Orbit.Intersection          76.2993
## Relative.Velocity.in.KM.per.sec     58.8524
## Perihelion.Distance                 58.3352
## Eccentricity                       53.1304
## Close.Approach.Date                 26.3045
## Aphelion.Dist                      21.3395
## Jupiter.Tisserand.Invariant          12.2465
## Miss.Dist.in.KM                    10.0205
## Mean.Anomaly                       10.0178
## Mean.Motion                        7.0710
## Semi.Major.Axis                    7.0710
## Orbital.Period                     7.0710
## Inclination                        3.4734
## Perihelion.Arg                     0.9156
## Asc.Node.Longitude                  0.0000
```

Analyze the data by creating a confusion matrix and generating performance statistics.

```
confusionMatrix(pihatcv.knn$pred, pihatcv.knn$obs, positive="True")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction False True
##      False  2551  272
##      True    93  162
##
##           Accuracy : 0.8814
##           95% CI : (0.8695, 0.8926)
##      No Information Rate : 0.859
##      P-Value [Acc > NIR] : 0.0001443
##
##           Kappa : 0.4085
##
##  Mcnemar's Test P-Value : < 2.2e-16
##
##           Sensitivity : 0.37327
##           Specificity : 0.96483
##      Pos Pred Value : 0.63529
##      Neg Pred Value : 0.90365
##           Prevalence : 0.14100
##      Detection Rate : 0.05263
##      Detection Prevalence : 0.08285
##      Balanced Accuracy : 0.66905
##
##      'Positive' Class : True
##
```

```
# Confusion matrix
conf.knn <- table(pihatcv.knn$pred, pihatcv.knn$obs)
conf.knn
```

```
##
##           False True
##      False  2551  272
##      True    93  162
```

```
knn.stats <- get_stats(conf.knn)
knn.stats
```

```
##           name      value
## 1      accuracy 0.8814165
## 2      error rate 0.1185835
## 3      precision 0.6352941
## 4      sensitivity 0.3732719
## 5      specificity 0.9648260
## 6      F-measure 0.4702467
## 7 Matthew's CC 0.4268670
```


Classification Tree Analysis

There are no model assumptions to be satisfied for decision trees. We are fitting a classification tree as opposed to a regression tree because the response variable is categorical and binary.

```
set.seed(1)
# rpart
nasa.CVrpart <- train(Hazardous ~ ., data=nasa,
                      method="rpart",
                      tuneGrid = expand.grid(cp=seq(0.005, 0.05, length=10)),
                      trControl=trainControl(method="cv", number=5,
                                              savePredictions=TRUE,
                                              classProbs=TRUE,
                                              selectionFunction = "oneSE"))

nasa.CVrpart

## CART
##
## 3078 samples
## 20 predictor
## 2 classes: 'False', 'True'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 2462, 2463, 2463, 2462, 2462
## Resampling results across tuning parameters:
##
##   cp      Accuracy   Kappa
## 0.005 0.9948026 0.978491
## 0.010 0.9948026 0.978543
## 0.015 0.9948026 0.978543
## 0.020 0.9948026 0.978543
## 0.025 0.9948026 0.978543
## 0.030 0.9948026 0.978543
## 0.035 0.9948026 0.978543
## 0.040 0.9948026 0.978543
## 0.045 0.9948026 0.978543
## 0.050 0.9948026 0.978543
##
## Accuracy was used to select the optimal model using the one SE rule.
## The final value used for the model was cp = 0.05.
# print tree
fancyRpartPlot(nasa.CVrpart$finalModel)
```



Rattle 2023–May–01 19:27:38 madelinepfister


If the plain `rpart` method is utilized, it selects a `cp` (complexity parameter) value of 0.05 to maximize accuracy for the final model. Note that the accuracy and kappa values are identical for each of the `cp` values tried by the model, so any choice within that range should produce comparable results.

We then assessed the predictive performance by plotting the ROC curve and finding the area under the curve.

```

pihatcv.rpart <- nasa.CVrpart$pred[nasa.CVrpart$pred$cp == 0.05,]
predcv.rpart <- prediction(pihatcv.rpart$True, pihatcv.rpart$obs)
perfcv.rpart <- performance(predcv.rpart, "tpr", "fpr")
plot(perfcv.rpart)

```



```
aucCV.rpart <- performance(predcv.rpart, "auc")@y.values
aucCV.rpart
```


```
## [[1]]
## [1] 0.9917045
```

The 10 most important variables in the model are:

```
# variable importance
varImp(nasa.CVrpart)
```

```
## rpart variable importance
##
##                                     Overall
## Minimum.Orbit.Intersection      100.000
## Est.Dia.in.KM.min.              17.973
## Absolute.Magnitude              17.973
## Est.Dia.in.KM.max.              17.973
## Est.Dia.in.KM.range             17.973
## Perihelion.Distance             12.601
## Miss.Dist.in.KM                 8.932
## Inclination                     6.388
## Relative.Velocity.in.KM.per.sec 2.229
## Jupiter.Tisserand.Invariant      0.000
## Orbital.Period                  0.000
## Mean.Motion                     0.000
## Mean.Anomaly                    0.000
## Perihelion.Arg                  0.000
## Close.Approach.Date             0.000
## Semi.Major.Axis                 0.000
## Asc.Node.Longitude              0.000
## Eccentricity                    0.000
## Aphelion.Dist                   0.000
## Orbit.Uncertainty               0.000
```

```
plot(varImp(nasa.CVrpart))
```



Analyze the data by creating a confusion matrix and generating performance statistics.

```
confusionMatrix(pihatcv.rpart$pred, pihatcv.rpart$obs, positive="True")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction False True
##      False 2637   9
##      True   7  425
##
##              Accuracy : 0.9948
##              95% CI : (0.9916, 0.997)
##      No Information Rate : 0.859
##      P-Value [Acc > NIR] : <2e-16
##
##              Kappa : 0.9785
##
##  McNemar's Test P-Value : 0.8026
##
##              Sensitivity : 0.9793
##              Specificity : 0.9974
##      Pos Pred Value : 0.9838
##      Neg Pred Value : 0.9966
##      Prevalence : 0.1410
##      Detection Rate : 0.1381
##      Detection Prevalence : 0.1404
##      Balanced Accuracy : 0.9883
##
##      'Positive' Class : True
```

```
##
# Confusion matrix
conf.rpart <- table(pihatcv.rpart$pred, pihatcv.rpart$obs)
conf.rpart

##
##           False True
## False  2637    9
## True     7   425

rpart.stats <- get_stats(conf.rpart)
rpart.stats

##           name          value
## 1    accuracy 0.994801819
## 2   error rate 0.005198181
## 3    precision 0.983796296
## 4  sensitivity 0.979262673
## 5  specificity 0.997352496
## 6    F-measure 0.981524249
## 7 Matthew's CC 0.978503227
```

Random Forest Analysis

There are no model assumptions to be satisfied for random forest analysis.

```
set.seed(1)
nasa.rf <- randomForest(as.factor(Hazardous) ~ ., data=nasa)
nasa.rf

##
## Call:
## randomForest(formula = as.factor(Hazardous) ~ ., data = nasa)
##               Type of random forest: classification
##               Number of trees: 500
## No. of variables tried at each split: 4
##
##               OOB estimate of  error rate: 0.49%
## Confusion matrix:
##      False True class.error
## False  2640    4 0.001512859
## True    11  423 0.025345622
```


```
set.seed(1)
nasa.CVrf <- train(Hazardous ~ ., data=nasa,
                  method="rf",
                  trControl=trainControl(method="cv", number=5,
                                          savePredictions=TRUE,
                                          classProbs=TRUE))
nasa.CVrf
```

```
## Random Forest
##
## 3078 samples
## 20 predictor
## 2 classes: 'False', 'True'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 2462, 2463, 2463, 2462, 2462
## Resampling results across tuning parameters:
##
##  mtry  Accuracy  Kappa
##    2    0.9944774 0.9769498
##   11    0.9954519 0.9811428
##   20    0.9961013 0.9838326
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 20.
```

The final model selects an mtry tuning parameter value of 20 in order to maximize accuracy.

We then assessed the model predictive performance by plotting the ROC curve and finding the model under the curve.

```
pihatcv.rf <- nasa.CVrf$pred[nasa.CVrf$pred$mtry == 20,]
predcv.rf <- prediction(pihatcv.rf$True, pihatcv.rf$obs)
perfcv.rf <- performance(predcv.rf, "tpr", "fpr")
plot(perfcv.rf)
```

```
aucCV.rf <- performance(predcv.rf, "auc")@y.values
aucCV.rf
```

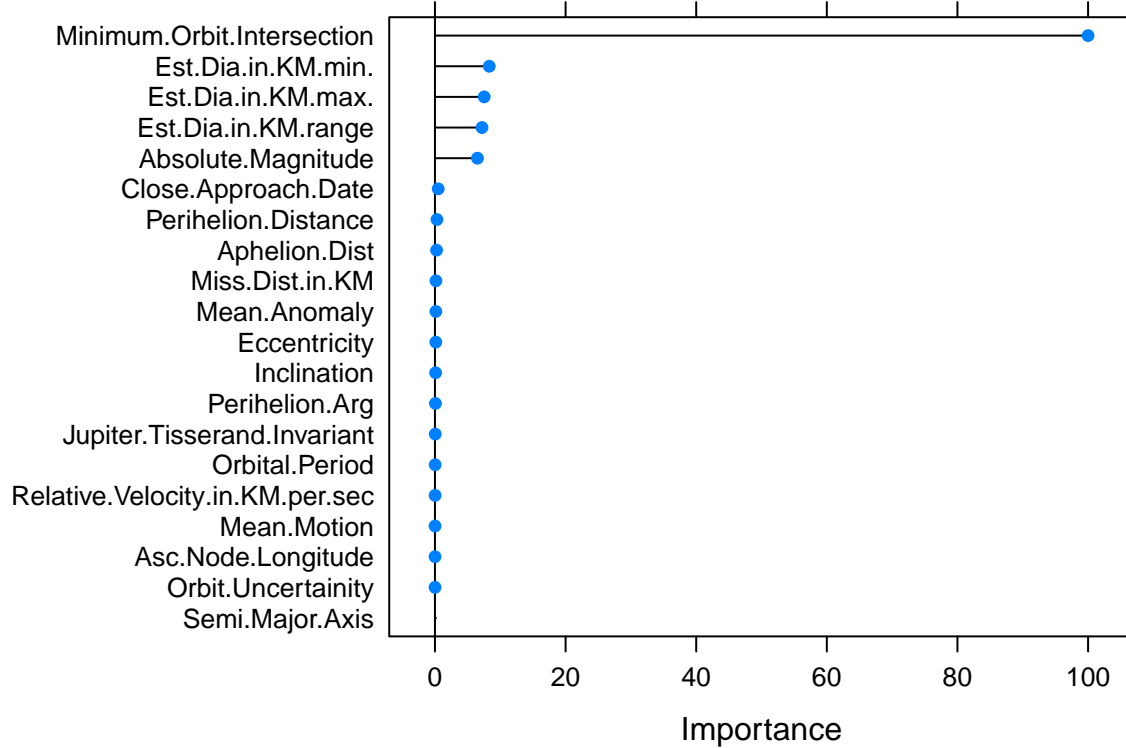
```
## [[1]]
## [1] 0.9986715
```

The 10 most important variables is the random forest model are:

```
varImp(nasa.CVrf)
```

```
## rf variable importance
##
##                               Overall
## Minimum.Orbit.Intersection    100.00000
## Est.Dia.in.KM.min.             8.31765
## Est.Dia.in.KM.max.             7.54595
## Est.Dia.in.KM.range           7.19967
## Absolute.Magnitude            6.50834
## Close.Approach.Date           0.50215
## Perihelion.Distance           0.31000
## Aphelion.Dist                 0.24640
## Miss.Dist.in.KM              0.14536
## Mean.Anomaly                 0.14268
## Eccentricity                 0.13382
## Inclination                  0.09681
## Perihelion.Arg               0.07256
## Jupiter.Tisserand.Invariant   0.04657
## Orbital.Period               0.03027
## Relative.Velocity.in.KM.per.sec 0.02670
## Mean.Motion                  0.01485
## Asc.Node.Longitude           0.01116
## Orbit.Uncertainty            0.01089
## Semi.Major.Axis              0.00000
```

```
plot(varImp(nasa.CVrf))
```



Analyze the data by creating a confusion matrix and generating performance statistics.

```
confusionMatrix(pihtcv.rf$pred, pihtcv.rf$obs, positive="True")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction False True
##      False  2640    8
##      True     4  426
##
##              Accuracy : 0.9961
##              95% CI : (0.9932, 0.998)
##      No Information Rate : 0.859
##      P-Value [Acc > NIR] : <2e-16
##
##              Kappa : 0.9838
##
##  Mcnemar's Test P-Value : 0.3865
##
##              Sensitivity : 0.9816
##              Specificity : 0.9985
##      Pos Pred Value : 0.9907
##      Neg Pred Value : 0.9970
##      Prevalence : 0.1410
##      Detection Rate : 0.1384
##      Detection Prevalence : 0.1397
##      Balanced Accuracy : 0.9900
```

```
##
##      'Positive' Class : True
##
# Confusion matrix
conf.rf <- table(pihatcv.rf$obs, pihatcv.rf$pred)
conf.rf

##
##      False True
## False  2640   4
##  True     8 426

rf.stats <- get_stats(conf.rf)
rf.stats

##      name      value
## 1  accuracy 0.996101365
## 2  error rate 0.003898635
## 3  precision 0.981566820
## 4  sensitivity 0.990697674
## 5  specificity 0.996978852
## 6   F-measure 0.986111111
## 7 Matthew's CC 0.983857862
```

Model Comparisons

Variable Importance Comparison

```
imp.knn <- rownames(varImp(fit.knn)$importance)
sel.knn <- imp.knn[order(varImp(fit.knn)$importance[,1], decreasing=T)][1:10]
imp.rp <- rownames(varImp(nasa.CVrpart)$importance)
sel.rp <- imp.rp[order(varImp(nasa.CVrpart)$importance[,1], decreasing=T)][1:10]
imp.rf <- rownames(varImp(nasa.CVrf)$importance)
sel.rf <- imp.rf[order(varImp(nasa.CVrf)$importance[,1], decreasing=T)][1:10]
intersect(sel.names, intersect(sel.knn, intersect(sel.rp, sel.rf)))
```

```
## [1] "Absolute.Magnitude"      "Est.Dia.in.KM.min."
## [3] "Minimum.Orbit.Intersection"
```

Predictive Performance Stat Comparison

```
df_merge <- merge(lasso.stats,knn.stats,by="name")
colnames(df_merge)[colnames(df_merge) == "value.x"] = "Lasso"
colnames(df_merge)[colnames(df_merge) == "value.y"] = "kNN"
df_merge <- merge(df_merge,rpart.stats,by="name")
df_merge <- merge(df_merge,rf.stats,by="name")
colnames(df_merge)[colnames(df_merge) == "value.x"] = "CT"
colnames(df_merge)[colnames(df_merge) == "value.y"] = "RF"
df_merge
```

##	name	Lasso	kNN	CT	RF
## 1	accuracy	0.94704353	0.8814165	0.994801819	0.996101365
## 2	error rate	0.05295647	0.1185835	0.005198181	0.003898635
## 3	F-measure	0.80710059	0.4702467	0.981524249	0.986111111
## 4	Matthew's CC	0.77682255	0.4268670	0.978503227	0.983857862
## 5	precision	0.82968370	0.6352941	0.983796296	0.981566820
## 6	sensitivity	0.78571429	0.3732719	0.979262673	0.990697674
## 7	specificity	0.97352496	0.9648260	0.997352496	0.996978852