

Analysis: Classification Trees & Random Forest

Hannah Norman

```
library(caret)

## Loading required package: ggplot2
## Loading required package: lattice
library(mlbench)
library(rattle)

## Loading required package: tibble
## Loading required package: bitops
## Rattle: A free graphical interface for data science with R.
## Version 5.5.1 Copyright (c) 2006-2021 Togaware Pty Ltd.
## Type 'rattle()' to shake, rattle, and roll your data.
library(randomForest)

## randomForest 4.7-1.1
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
## The following object is masked from 'package:rattle':
##
##     importance
## The following object is masked from 'package:ggplot2':
##
##     margin
library(ROCR)
```

Original NASA Asteroids Data Set

```
nasa_orig <- read.csv("nasa_original.csv")
# nasa_orig
```

Cleaned NASA Asteroids Data Set

```
nasa <- read.csv("nasa.csv")
#nasa
```

Numerical and Graphical Summaries of Response Variable


```
prop.hazardous <- prop.table(table(nasa$Hazardous))
prop.hazardous
```

```
##
##      False      True
## 0.8589994 0.1410006
```

```
par(mfrow=c(1,2))
```

```
# pie chart
count.hazardous <- table(nasa$Hazardous)
lbls <- paste(levels(as.factor(nasa$Hazardous)), ":", " ",
              round(prop.hazardous,3)*100, "%", sep=" ")
pie(count.hazardous, labels=lbls, col=c("salmon", "palegreen"))
```

```
# bar plot
barplot(prop.hazardous, xlab="Hazardous", ylab="Proportion", ylim=c(0, 1.0),
        col=c("salmon", "palegreen"))
```



Analyses of Predictors

(in search of outliers that might skew analysis)

```
# numerical summary + box plots
summary(nasa)
```

```
## Absolute.Magnitude Est.Dia.in.KM.min. Est.Dia.in.KM.max. Est.Dia.in.KM.range
## Min. :15.20 Min. :0.00101 Min. :0.00226 Min. :0.001249
## 1st Qu.:20.40 1st Qu.:0.03052 1st Qu.:0.06824 1st Qu.:0.037722
## Median :22.30 Median :0.09216 Median :0.20608 Median :0.113919
## Mean :22.53 Mean :0.17185 Mean :0.38428 Mean :0.212423
```

```
## 3rd Qu.:24.70      3rd Qu.:0.22108      3rd Qu.:0.49436      3rd Qu.:0.273273
## Max. :32.10      Max. :2.42412      Max. :5.42051      Max. :2.996383
## Close.Approach.Date Relative.Velocity.in.KM.per.sec Miss.Dist.in.KM
## Min. :19950101      Min. : 0.8002      Min. : 26610
## 1st Qu.:20010808      1st Qu.: 8.1680      1st Qu.:16215102
## Median :20070912      Median :12.3870      Median :37033462
## Mean :20066259      Mean :13.6128      Mean :36473114
## 3rd Qu.:20120922      3rd Qu.:17.5079      3rd Qu.:56290664
## Max. :20160908      Max. :43.7899      Max. :74781600
## Orbit.Uncertainty Minimum.Orbit.Intersection Jupiter.Tisserand.Invariant
## Min. :0.0      Min. :0.0000021      Min. :2.196
## 1st Qu.:1.0      1st Qu.:0.0151341      1st Qu.:3.807
## Median :5.0      Median :0.0473452      Median :4.798
## Mean :4.1      Mean :0.0823254      Mean :4.852
## 3rd Qu.:7.0      3rd Qu.:0.1249268      3rd Qu.:5.774
## Max. :9.0      Max. :0.4778910      Max. :9.025
## Eccentricity Semi.Major.Axis Inclination Asc.Node.Longitude
## Min. :0.01296      Min. :0.6159      Min. : 0.01451      Min. : 0.0019
## 1st Qu.:0.24907      1st Qu.:1.0530      1st Qu.: 4.78566      1st Qu.: 83.6648
## Median :0.38236      Median :1.3347      Median : 9.68715      Median :173.5898
## Mean :0.39321      Mean :1.4848      Mean :12.84105      Mean :173.5232
## 3rd Qu.:0.52589      3rd Qu.:1.8386      3rd Qu.:18.38238      3rd Qu.:258.4855
## Max. :0.96026      Max. :3.9908      Max. :75.40667      Max. :359.9059
## Orbital.Period Perihelion.Distance Perihelion.Arg Aphelion.Dist
## Min. : 176.6      Min. :0.08074      Min. : 0.0069      Min. :0.8038
## 1st Qu.: 394.7      1st Qu.:0.67807      1st Qu.: 95.6381      1st Qu.:1.3191
## Median : 563.2      Median :0.87434      Median :188.4906      Median :1.7911
## Mean : 692.8      Mean :0.84323      Mean :184.1273      Mean :2.1264
## 3rd Qu.: 910.6      3rd Qu.:1.01936      3rd Qu.:272.5434      3rd Qu.:2.7517
## Max. :2912.0      Max. :1.29983      Max. :359.9931      Max. :6.8918
## Mean.Anomaly Mean.Motion Hazardous
## Min. : 0.0032      Min. :0.1236      Length:3078
## 1st Qu.: 83.3164      1st Qu.:0.3953      Class :character
## Median :183.8903      Median :0.6392      Mode :character
## Mean :180.1357      Mean :0.6811
## 3rd Qu.:276.3132      3rd Qu.:0.9121
## Max. :359.9180      Max. :2.0390
```


```
par(mfrow=c(2,4))
boxplot(nasa$Est.Dia.in.KM.range, main="Est.Dia.in.KM.range")
boxplot(nasa$Relative.Velocity, main="Relative.Velocity")
boxplot(nasa$Minimum.Orbit.Intersection, main="Minimum.Orbit.Intersection")
boxplot(nasa$Inclination, main="Inclination")
boxplot(nasa$Orbital.Period, main="Orbital.Period")
boxplot(nasa$Aphelion.Dist, main="Aphelion.Dist")
boxplot(nasa$Mean.Motion, main="Mean.Motion")
```

```
# to find row number of outlier observations
nasa[which.max(nasa$Est.Dia.in.KM.range),]
```


```
## Absolute.Magnitude Est.Dia.in.KM.min. Est.Dia.in.KM.max.
## 2728 15.2 2.424125 5.420508
## Est.Dia.in.KM.range Close.Approach.Date Relative.Velocity.in.KM.per.sec
## 2728 2.996383 20141222 23.5172
## Miss.Dist.in.KM Orbit.Uncertainty Minimum.Orbit.Intersection
```

```
## 2728      45467472      0      0.153116
##      Jupiter.Tisserand.Invariant Eccentricity Semi.Major.Axis Inclination
## 2728      4.864      0.3462176      1.261473      36.90474
##      Asc.Node.Longitude Orbital.Period Perihelion.Distance Perihelion.Arg
## 2728      111.2844      517.5058      0.8247287      349.1364
##      Aphelion.Dist Mean.Anomaly Mean.Motion Hazardous
## 2728      1.698217      328.5237      0.6956443      False
```


Est.Dia.in.KM.range




Relative.Velocity




Minimum.Orbit.Intersect




Inclination




Orbital.Period



Aphelion.Dist



Mean.Motion



Predictive Performance Stats

```
get_stats <- function(CM) {
  TP <- CM[2,2]
  FP <- CM[1,2]
  TN <- CM[1,1]
  FN <- CM[2,1]

  acc <- (TP+TN) / (TP+TN+FN+FP)
  err <- (FP+FN) / (TP+TN+FN+FP)
  pre <- (TP) / (TP+FP)
  sen <- (TP) / (TP+FN)
  spe <- (TN) / (TN+FP)
  fme <- (2*pre*sen) / (pre+sen)
  mcc_denom <- sqrt(TP+FP)*sqrt(TP+FN)*sqrt(TN+FP)*sqrt(TN+FN)
  mcc <- (TP*TN - FP*FN) / mcc_denom

  name <- c("accuracy", "error rate", "precision", "sensitivity", "specificity", "F-measure", "Matthew's")
  value <- c(acc, err, pre, sen, spe, fme, mcc)
  stats <- data.frame(name, value)

  return (stats)
}
```

Classification Tree Analysis

1. Make sure that the model assumptions, if any, are satisfied.

No model assumptions of decision trees to be satisfied? We are fitting a classification tree as opposed to a regression tree because the response variable is categorical and binary.

2. Assess the model fit and perform diagnostics, if appropriate.

Both methods appear to give identical results. Note the identical accuracy and kappa ratings across all cp tuning parameters in the plain rpart method...


```
set.seed(1)

# rpart
nasa.CVrpart <- train(Hazardous ~ ., data=nasa,
                      method="rpart",
                      tuneGrid = expand.grid(cp=seq(0.005, 0.05, length=10)),
                      trControl=trainControl(method="cv", number=10,
                                              savePredictions=TRUE,
                                              classProbs=TRUE,
                                              selectionFunction = "oneSE"))

nasa.CVrpart

## CART
##
## 3078 samples
## 20 predictor
## 2 classes: 'False', 'True'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 2770, 2771, 2771, 2769, 2770, 2771, ...
## Resampling results across tuning parameters:
##
##   cp      Accuracy   Kappa
## 0.005 0.9947999 0.9783443
## 0.010 0.9947999 0.9783443
## 0.015 0.9947999 0.9783443
## 0.020 0.9947999 0.9783443
## 0.025 0.9947999 0.9783443
## 0.030 0.9947999 0.9783443
## 0.035 0.9947999 0.9783443
## 0.040 0.9947999 0.9783443
## 0.045 0.9947999 0.9783443
## 0.050 0.9947999 0.9783443
##
## Accuracy was used to select the optimal model using the one SE rule.
## The final value used for the model was cp = 0.05.

# print tree
fancyRpartPlot(nasa.CVrpart$finalModel)
```



Rattle 2023-Apr-26 23:40:22 hannah

3. Identify tuning parameters to be used, if appropriate.


If the plain `rpart` method is utilized, it selects a `cp` (complexity parameter) value of 0.05 to maximize accuracy for the final model. Note that the accuracy and kappa values are identical for each of the `cp` values tried by the model, so any choice within that range should produce comparable results.

4. Identify and interpret the effect of selected variables.

```
# variable importance (rpart)
varImp(nasa.CVrpart)
```

```
## rpart variable importance
##
##                               Overall
## Minimum.Orbit.Intersection    100.000
## Est.Dia.in.KM.range           17.973
## Est.Dia.in.KM.min             17.973
## Absolute.Magnitude            17.973
## Est.Dia.in.KM.max             17.973
## Perihelion.Distance           12.601
## Miss.Dist.in.KM               8.932
## Inclination                   6.388
## Relative.Velocity.in.KM.per.sec 2.229
## Aphelion.Dist                 0.000
## Perihelion.Arg                0.000
## Jupiter.Tisserand.Invariant   0.000
## Mean.Motion                   0.000
## Orbit.Uncertainty             0.000
## Close.Approach.Date           0.000
## Mean.Anomaly                  0.000
## Semi.Major.Axis               0.000
## Asc.Node.Longitude            0.000
## Orbital.Period                0.000
## Eccentricity                  0.000
```

```
plot(varImp(nasa.CVrpart))
```




5. Evaluate the cross-validated (CV) predictive performance.

```
head(nasa.CVrpart$pred)
```

```
##      pred  obs rowIndex      False      True      cp Resample
## 1 False False      11 1.000000000 0.000000000 0.005  Fold01
## 2 False False      13 0.993243243 0.006756757 0.005  Fold01
## 3 False False      34 1.000000000 0.000000000 0.005  Fold01
## 4 False False      38 0.993243243 0.006756757 0.005  Fold01
## 5  True  True      68 0.002624672 0.997375328 0.005  Fold01
## 6  True  True      70 0.002624672 0.997375328 0.005  Fold01
```

```
pihatcv.rpart <- nasa.CVrpart$pred[nasa.CVrpart$pred$cp == 0.05,]
predcv.rpart <- prediction(pihaticv.rpart$True, pihaticv.rpart$obs)
perfcv.rpart <- performance(predcv.rpart, "tpr", "fpr")
plot(perfcv.rpart)
abline(a=0, b=1, lty=2)
```



```
aucCV.rpart <- performance(predcv.rpart, "auc")@y.values
aucCV.rpart
```

```
## [[1]]
## [1] 0.991193
```

```
confMat <- table(pihatcv.rpart$obs, pihatcv.rpart$pred)
confMat
```

```
##
##           False True
## False  2640     4
## True    12    422
```

```
rpart.stats <- get_stats(confMat)
rpart.stats
```

```
##           name      value
## 1      accuracy 0.994801819
## 2      error rate 0.005198181
## 3      precision 0.990610329
## 4      sensitivity 0.972350230
## 5      specificity 0.998487141
## 6      F-measure 0.981395349
## 7 Matthew's CC 0.978431703
```

Random Forest Analysis

1. Make sure that the model assumptions, if any, are satisfied.

No model assumptions of random forest to be satisfied?

2. Assess the model fit and perform diagnostics, if appropriate.

```
set.seed(1)
```



```
nasa.rf <- randomForest(as.factor(Hazardous) ~ ., data=nasa)
nasa.rf

##
## Call:
## randomForest(formula = as.factor(Hazardous) ~ ., data = nasa)
##           Type of random forest: classification
##           Number of trees: 500
## No. of variables tried at each split: 4
##
##           OOB estimate of  error rate: 0.49%
## Confusion matrix:
##           False True class.error
## False  2640    4 0.001512859
## True    11  423 0.025345622
```

```
set.seed(1)

nasa.CVrf <- train(Hazardous ~ ., data=nasa,
                  method="rf",
                  trControl=trainControl(method="cv", number=10,
                                          savePredictions=TRUE,
                                          classProbs=TRUE))
nasa.CVrf
```

```
## Random Forest
##
## 3078 samples
## 20 predictor
## 2 classes: 'False', 'True'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 2770, 2771, 2771, 2769, 2770, 2771, ...
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa
## 2 0.9951256 0.9796031
## 11 0.9960986 0.9837768
## 20 0.9964264 0.9850905
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 20.
```

3. Identify tuning parameters to be used, if appropriate.

The final model selects an mtry tuning parameter value of 20 in order to maximize accuracy.


4. Identify and interpret the effect of selected variables.

```
varImp(nasa.CVrf)
```

```
## rf variable importance
##
##           Overall
## Minimum.Orbit.Intersection 100.00000
## Est.Dia.in.KM.range       7.72334
```

```
## Est.Dia.in.KM.max.          7.60017
## Absolute.Magnitude          7.29260
## Est.Dia.in.KM.min.          6.23992
## Close.Approach.Date         0.51729
## Perihelion.Distance         0.31224
## Aphelion.Dist               0.22185
## Miss.Dist.in.KM            0.15809
## Eccentricity                0.15804
## Mean.Anomaly                0.14305
## Inclination                 0.09922
## Perihelion.Arg              0.07231
## Jupiter.Tisserand.Invariant  0.05422
## Asc.Node.Longitude          0.04687
## Relative.Velocity.in.KM.per.sec 0.04404
## Mean.Motion                 0.03637
## Orbital.Period              0.03556
## Semi.Major.Axis             0.02316
## Orbit.Uncertainty           0.00000
```

```
plot(varImp(nasa.CVrf))
```




5. Evaluate the cross-validated (CV) predictive performance.

```
head(nasa.CVrf$pred)
```

```
##   pred  obs False  True rowIndex mtry Resample
## 1 False False 0.970 0.030      11    2  Fold01
## 2 False False 0.984 0.016      13    2  Fold01
## 3 False False 0.968 0.032      34    2  Fold01
## 4 False False 0.996 0.004      38    2  Fold01
## 5  True  True 0.242 0.758      68    2  Fold01
```

```
## 6 True True 0.212 0.788      70    2 Fold01
pihatcv.rf <- nasa.CVrf$pred[nasa.CVrf$pred$mtry == 20,]
predcv.rf <- prediction(pihatcv.rf$True, pihatcv.rf$obs)
perfcv.rf <- performance(predcv.rf, "tpr", "fpr")
plot(perfcv.rf)
abline(a=0, b=1, lty=2)
```



```
aucCV.rf <- performance(predcv.rf, "auc")@y.values
aucCV.rf
```

```
## [[1]]
## [1] 0.9986688
```

```
confMat <- table(pihatcv.rf$obs, pihatcv.rf$pred)
confMat
```

```
##
##      False True
## False 2642    2
## True   9    425
```

```
rf.stats <- get_stats(confMat)
rf.stats
```

```
##      name      value
## 1 accuracy 0.996426251
## 2 error rate 0.003573749
## 3 precision 0.995316159
## 4 sensitivity 0.979262673
## 5 specificity 0.999243570
## 6 F-measure 0.987224158
## 7 Matthew's CC 0.985190895
```